Group table pack

Group: ______ *(a short, friendly name that does not identify anyone — e.g. "Otters", "Team Kelp")* · **Track:** ______ (A · B · C · D)

Everything your group needs for the applied session is in this booklet. You have about **fifteen minutes** in Part 2. Get the **core three** done first — and if that is all you manage, you have done the task:

1. **One artefact move** — run your tool on your problem and make something: a critique, a template, a small artefact, a translation.
2. **One or two caught errors** — moments it was fluent and wrong. Keep them.
3. **One insight** — the limitation, the safeguard, or one honest thing this exposed about how your field is over- or under-using AI, that you will give in the 45-second round.

If you have about five minutes left, go deeper:

4. **The automation–steering map** — which steps the tool ran, and where *you* steered.
5. **Interwoven or staged?** — was your oversight woven through every step, or staged at checkpoints between phases? Say why, and what it cost.

A rough split that keeps you on time: **~6 min make · ~3 min break it · ~2 min insight · ~4 min map & oversight if it is left.** When you fall behind, drop the map before the insight.

The full rubric scoring and the societal reflection come **later**, when you tidy and share the note — not in the fifteen minutes.

Before you paste anything — the red lines

Never put into a free AI tool: personal data; students' marks, references or pastoral notes; unpublished or participant data; or confidential, peer-review or embargoed material. When in doubt, **anonymise, synthesise or abstain**. The one-page data decision aid in this pack walks you through it.

Your shared note

Your group's note lives in the **workshop app** — scan the QR or go to **genai-rt.web.app**, and your facilitator reads out the **session passcode**. One device starts the group; the others join with the code it shows. Fill the essentials as you go and **submit by 13:45**. No account, nothing to install.

> **One note, one surface.** Use the **app** (the default). If you have no device or it will not connect, **HackMD** (hackmd.io) is the fallback — same headings. Not both.

One reflection for the wrap-up

Not for the fifteen minutes — save it for when you tidy the note, or use it as your lightning-round line if it is the sharper one: **is your field generally under- or over-using GenAI for this kind of task?** Why, with what consequences — and how would you redress it over the next two years?

What is in this pack

- **Role cards** — cut along the lines and take one each (five angles, one per person).
- **Data decision aid** — the "can I paste this?" one-pager.
- **The rubric** — what you are working towards (you score it when you tidy the note).
- **A worked example per track** — your default problem if no one brought a real one.
- **Starter prompts** — copy, adapt, go.

A personal selection, made in the facilitator's own capacity. It is not the position of the University of Oxford (the facilitator's employer) or of the host, and it is not legal advice.

Group role cards (optional prompt sheet)

Five angles to keep an eye on in a group of five — take one each, or divide them however suits you. They are about **where attention goes**, not who owns the work; rotate the keyboard freely. Each maps to one **rubric dimension**, so make sure none is left unwatched. **An experienced group can ignore these entirely** — just don't let the building crowd out the interrogating.

> **Print tip.** Optional. Print one set per group if it helps; cut along the rules. At four people, the Convenor doubles as Steward. At six, add a second Sceptic — more red-teaming is no bad thing.

The mapping (canonical):

Angle (handle)	Rubric dimension	In one line
Convenor	Project Definition	Keeps the real problem and the clock in focus
Driver	Technology Stack	Operates the tools, owns the tool choices
Steward	Data Security & Ethics	Guards the red lines
Reporter	Financial & Scalability	Owns the record, the submission and the 45-second insight
Sceptic	Human-in-the-Loop Protocol	Tries to break the output, curates caught errors

---

✂ —

CONVENOR — *Project Definition*

Your angle: the group stays on a real problem and on time.

- **Do:** in the first minute, agree your **one** problem — the seed you formed around, or a real one a member brought — and a success criterion; split the five angles; watch the clock against the timeline; make sure everyone speaks.
- **At the pivot (~13:26):** call the pause; lead the four mid-point questions; move the group from *building* to *interrogating*.
- **Ask often:** "What is the actual problem, and how will we know it helped?"

✂ —

DRIVER — *Technology Stack*

Your angle: operating the tools and the tech choices.

- **Do:** share your screen; place your tool on the spectrum (off-the-shelf / no-code / IDE–API); be ready to say why this level and not one up or down.
- **At the pivot (~13:26):** note what the tool made genuinely easier — and the cost.

- **Ask often:** "Would a different level of tool change what we can trust?"

✂ —

STEWARD — *Data Security & Ethics*

Your angle: holding the red lines.

- **Do:** vet everything *before* it is pasted (use the data decision aid); watch for personal, special-category or confidential data; track disclosure, IP and consent.
- **At the pivot (~13:26):** confirm nothing crossed a red line; flag anything to redact.
- **Ask often:** "Whose data is this, and would we be content for it to be public?"

✂ —

REPORTER — *Financial & Scalability*

Your angle: the record and the insight.

- **Do:** keep the group's note in the app, in the rubric template; judge whether the approach survives the free tier and scale; submit the note for review during Part 2 (by **13:45**); give the **45-second** insight at 13:33 (no slides).
- **At the pivot (~13:26):** make sure the rubric is being filled in, not left to the end.
- **Ask often:** "Does this still work next month, for fifty people, for free?"

✂ —

SCEPTIC — *Human-in-the-Loop Protocol*

Your angle: breaking the output, and deciding where the human must stay.

- **Do:** stress-test every result; hunt hallucinations, bias and false confidence; curate the **museum of caught errors**; name the checkpoints and what you will never delegate; draw the **automation–steering map** and decide whether oversight is interwoven or staged.
- **At the pivot (~13:26):** bring your best caught error to the group.
- **Ask often:** "Where, exactly, does a human have to check before we trust this?"

✂ —

A personal selection, made in the facilitator's own capacity. It is not the position of the University of Oxford (the facilitator's employer) or of the host, and it is not legal advice.

One-page data decision aid — *can I paste this in?*

Print one per group, or keep it open in a tab. Use it **before** anything goes into a free, third-party AI tool. It is a thinking aid, not legal advice; institutional policy, the UK GDPR and research ethics still govern.

The one-line test

> **If you would not pin it to a public noticeboard in the building, do not paste it into a free consumer AI tool.**

Free tiers commonly **reuse your input to train models**, and settings change without notice. Treat anything you type as potentially **non-private and non-retractable**.

Walk the four questions

1. Is it personal data? — does it relate to an identifiable living person (names, emails, student or participant records, anything that could single someone out, even indirectly)? → If **yes**, do not paste it as-is. Anonymise first, and check your lawful basis.

2. Is it special-category or especially sensitive? — health, ethnicity, religion, politics, sexuality, biometrics; or children's data; or commercially sensitive material. → If **yes**, **stop.** A free consumer tool is the wrong place. Anonymise beyond recognition, synthesise or abstain.

3. Is it confidential, unpublished or restricted? — unpublished data or manuscripts, grant drafts, peer-review files, NDA-covered material, anything under **embargo** or shared without co-author or participant **consent**. → If **yes**, **stop.** Confidentiality and embargoes are not yours alone to waive.

4. Do you hold the rights? — is the text or image yours to share, or is it someone else's intellectual property? → If unsure, treat it as restricted.

If you cleared all four, you are probably fine. If any gave you pause, that pause is the signal — act on it.

Green / amber / red, at a glance

🟢 Generally fine	🟠 Pause and treat with care	🔴 Do not paste into a free tool
Already-public text (your published abstract)	Lightly disguised real material	Personal data of identifiable people
Fully synthetic or fabricated examples	Aggregated or partially de-identified data	Special-category / children's data
Generic, non-identifying questions	Draft prose with no sensitive content	Unpublished data, grants, peer review
Your own writing you hold rights to	Anything you would hesitate to email widely	Confidential, NDA or embargoed material

---

If the answer is "no" — what to do instead

- **Anonymise** — remove names, identifiers and anything that singles someone out; check it cannot be re-identified by combination.
- **Synthesise** — recreate the *structure* of the problem with invented content. A critique of a method works just as well on a paraphrased method.
- **Abstract up a level** — discuss the shape of the task, not the sensitive specifics.
- **Use an appropriately-governed tool** — an enterprise or paid service under a proper data-processing agreement may be acceptable where a free consumer tool is not. That is a decision for your institution, not for today.
- **Abstain** — sometimes the right answer is not to use the tool. That is a valid finding, and a good one for your museum of caught errors.

Evaluation rubric & note template

This is the group's workbench for the day. By default you fill it in the **workshop app** (genai-rt.web.app) — it has these fields built in, so there is nothing to set up beyond the passcode. The **HackMD** steps below are the **fallback** if you cannot use the app. Either way it does two jobs:

1. **Hold your shared note** — in the app (default) or HackMD (fallback); no prior experience needed, nothing to install.
2. **Structure your thinking** around five dimensions, each scored **1–5** (**1 = Nascent · 3 = Developing · 5 = Robust**), plus a deliverable link, a museum of caught errors and your lightning-round insight.

You are not graded. The rubric is a **thinking tool and a shared record** — and, once archived, part of the open account of what this cohort learned.

> **Short-format note.** Part 2's applied session is about 15 minutes. Get the **core three** first — the problem and the tool, **one caught error**, and **one insight** for the lightning round. If you have about five minutes left, add the **automation–steering map** (interwoven or staged) and the societal reflection. A thoughtful partial note beats a rushed full one.

A. The HackMD fallback (≈ 2 minutes, Reporter drives — only if you cannot use the app)

1. Go to **hackmd.io** and sign in — it is **free**, and you can use a Google, GitHub or email login. There is **nothing to install**.
2. Click **"+ New note"**. The left pane is Markdown; the right pane shows a live preview. You can ignore the formatting and just type into the headings.
3. **Copy the whole template** in section B below (everything inside the box) and **paste it into the note**.
4. **Fill it in as you work**, not all at the end. Rough notes are fine.
5. **Make it link-shareable:** open the note's sharing/permissions control and set it so **anyone with the link can read** (you do *not* need to grant editing).
6. **Copy the note's link.**
7. **During Part 2 (by 13:45)**, share that link with the **facilitator** — show the note on screen with your **group number** and **track**. You do *not* need a GitHub account; the facilitator gathers every group's link.

> **If HackMD will not let you sign up** (some logins are restricted), do not lose time: write the same template in any shared document or on paper. It is plain Markdown and the facilitator can transcribe one note into the archive afterwards.

The facilitator exports every note to Markdown after the session and commits it under `submissions/`, so your work is preserved exactly.

B. The note template — the app has these fields built in; paste this box

into HackMD on the fallback

```
# [Group NN] · Track [A / B / C / D] · 24 June 2026

**Members & roles:** Convenor — … · Reporter — … · Driver — … · Sceptic — … · Steward — …
**The real problem we brought:** …

## 1. Project Definition
- The real problem, in one sentence a colleague would recognise: …
- Success criterion — how we would know the tool actually helped: …

## 2. Technology Stack
- Tool(s) used, and where they sit on the spectrum (off-the-shelf / no-code / IDE–API): …
- Why this level and not one up or down (control vs data exposure vs effort): …

## 3. Data Security & Ethics
- What we put into the tool, and how we de-identified it: …
- Red lines we held; UK GDPR, ethics, disclosure of AI use, IP, consent and fairness: …

## 4. Financial & Scalability Constraints
- Free-tier limits we hit or foresee: …
- What breaks first at scale — cost, usage limits or trust: …

## 5. Human-in-the-Loop Protocol
- **Automation–steering map** — break the project into its phases; for each, note what is *automated* (the tool does) and what is *human-steered* (we decide, verify or override):

| Phase | Automated (AI does) | Human steering (we decide / verify / override) |
|---|---|---|
| … | … | … |

- **Interwoven or staged?** Is our oversight **interwoven** (continuous — a human in the loop at every step) or **staged** (concentrated at distinct checkpoints between phases)? Which did we choose, and why?
- What our choice costs, and how we offset it. (Interwoven oversight is thorough but heavy, and checking every step can dull attention until you wave things through out of habit; staged oversight is lighter, but errors can pile up unnoticed in the stretches between checkpoints.): …
- Checkpoints we set *in advance*: …
- What we will never delegate to the tool: …
- Who is accountable for the final judgement: …

## Deliverable
- Link or screenshot of the artefact we made: …

## Museum of caught errors
| What happened | How we caught it | What it signals about where humans must stay |
|---|---|---|
| … | … | … |

## Societal reflection (complete when you tidy & share the note — no score, just think)
- Thinking vs writing — if the tool did the writing, what thinking did we still have to do, and what might we lose by offloading it? …
- Fairness — could this widen or narrow disparities (first vs second language, more vs less support, career stage, socioeconomic background)? Who gains, who is left behind? …
```

```
- Disclosure — would we disclose using AI here? Would we, or should we, disclose the conditions it offset
(e.g. writing in a second language, limited support)? Why or why not? …
- Disciplinary norms — is our field generally **under-** or **over-using** GenAI for this kind of task?
Why, with what consequences, and how would we redress it over the next two years? …

## Lightning-round insight — 45 seconds, no slides
> Pick the sharpest of: the single most significant **limitation** we found; the most
> important **human-in-the-loop safeguard** we built in; or one honest thing this exposed
> about how our **field** is over- or under-using AI for this task:
>
> …

## Self-assessment — complete when you tidy & share the note (1 = Nascent · 3 = Developing · 5 = Robust)
| Dimension | Score (1–5) | One-line justification |
|---|---|---|
| Project Definition | | |
| Technology Stack | | |
| Data Security & Ethics | | |
| Financial & Scalability | | |
| Human-in-the-Loop | | |
```

C. The rubric anchors

Use these to place yourselves honestly. Most groups will sit around **3** on most dimensions in a short applied session — that is exactly right. A thoughtful **3** with a sharp caught error beats an unexamined **5**.

1 · Project Definition

Is the real problem and its success criterion clear?

- **1 — Nascent:** The task is vague or a toy ("use AI on something"); no success criterion.
- **3 — Developing:** A real problem is named and there is a rough sense of what "better" would look like; scope may be broad.
- **5 — Robust:** A specific, bounded, genuinely-held problem a colleague would recognise, with an explicit success criterion you could actually test.

2 · Technology Stack

Right tool on the spectrum, and why?

- **1 — Nascent:** One tool used by default; no awareness of alternatives or of the spectrum.
- **3 — Developing:** The choice is explained, with some sense of where it sits and what moving a level up or down would change.
- **5 — Robust:** Deliberate placement on the spectrum, with a reasoned trade-off of control, data exposure and effort, and alternatives weighed.

3 · Data Security & Ethics

Red lines held; UK GDPR, ethics and fairness addressed?

- **1 — Nascent:** Personal, special-category or confidential data used without thought; no disclosure, consent or fairness consideration.
- **3 — Developing:** Red lines mostly respected and data de-identified; some ethical reflection on disclosure, IP, bias or fairness.
- **5 — Robust:** Red lines explicitly held; anonymisation or lawful basis reasoned; disclosure, IP, consent, bias and fairness addressed; the note would withstand being public.

4 · Financial & Scalability Constraints

Does the free tier hold; what breaks at scale?

- **1 — Nascent:** No thought to cost or scale; assumes free forever.
- **3 — Developing:** Free-tier limits noted, with a rough sense of what scaling would cost or require.
- **5 — Robust:** Clear-eyed on free-tier viability, lock-in and the precise point at which data, cost or trust breaks the approach.

5 · Human-in-the-Loop Protocol

Explicit checkpoints, verification, accountability — and a deliberate map of automation against human steering?

- **1 — Nascent:** Output trusted as-is; no checkpoint; no named accountability; no sense of where automation ends and human steering begins.
- **3 — Developing:** Some verification done; one or two checkpoints identified; an implicit split between the automated and human-steered steps.
- **5 — Robust:** A deliberate automation–steering map; oversight consciously chosen as interwoven or staged, with the weakness of that choice offset; what is never delegated is explicit; a named person owns the final judgement on correctness, ethics and attribution.

Disclaimer

These materials and the tools they reference are the facilitator's own choices, made in a personal capacity. They do not represent the views or official position of the University of Oxford, the facilitator's employer, or of the host, the University of Westminster. The University of Oxford accepts no liability for the selection, use or outcomes of any third-party tool. Participants remain solely responsible for compliance with their own institutional policy, the UK GDPR and research ethics.

This wording is a template and is not legal advice.

Worked examples — one per track

A worked example is the default for every group — use it unless a member actively brings a real, non-confidential problem they would rather work on. The material below is **fully synthetic and safe to paste** (no real people, data or unpublished work), which removes the de-identification deliberation that would otherwise eat into the fifteen minutes; each example is sized for the short Part 2 applied session and leads naturally to a caught error. It is also the fallback when the Wi-Fi fails and you switch to critiquing on paper.

Track A · Methodological Blind-Spot Detector

Scenario. You are about to submit the study below. Use an AI tool as a critical reviewer, then verify which of its objections are real, which are generic and which are wrong.

> *Synthetic methods snippet (safe to paste):* "We ran a cross-sectional online survey of 200 marketing managers recruited through a LinkedIn post. Each manager completed a 12-item self-report scale of 'authentic leadership' and, in the same questionnaire, rated their own team's 'innovativeness' on a 5-item scale. Authentic leadership correlated positively with team innovativeness ($r = .46$, $p < .001$). We conclude that authentic leadership **causes** higher team innovation and recommend leadership training to raise innovation."

First move. Ask the tool to review across explicit lenses: sampling and recruitment, common-method bias, construct validity, causal inference and generalisability. Then **classify each point** it raises as real / generic / wrong.

What to watch for. A strong tool should flag the **causal claim from cross-sectional data**, **common-method / single-source bias** (same person rates both variables), **self-selection** via LinkedIn and **self-report of one's own team**. Note whether it also pads the list with vague, always-true objections — and whether it ever invents a statistic or a citation. Those are your museum pieces.

Track B · Accessible Executive-Function Layer

Scenario. Turn a messy set of meeting notes into a clean, owned action list — then check the tool did not quietly invent or misattribute anything.

> *Synthetic meeting notes (safe to paste; roles only, no real people):* "Centre meeting. Talked about the seminar series — someone should sort speakers for the autumn, maybe the deputy director? Budget underspend, need to use it before July or we lose it. The new PhD reps raised workload worries. Website is out of date (old staff list). Agreed to revisit the mentoring scheme but no one said who. Ethics turnaround is too slow — chase the committee. Next meeting in three weeks, same room hopefully."

First move. Ask for a structured action list with **owner**, **deadline** and **priority**, plus a separate "unclear / needs an owner" list. Build it as a **reusable template**, not a one-off.

What to watch for. Does the tool **invent owners or dates** that the notes never stated ("Deputy Director — by 30 June")? Does it silently drop the items with no owner instead of flagging them? Does it confidently over-formalise a vague discussion? The safeguard you design — "every action must trace to a line in the notes; no invented owners" — is the deliverable.

Track C · Rapid Prototyping for Knowledge Translation *(technical stretch)*

Scenario. Build a tiny interactive that communicates a finding — and verify it is actually correct before you would ever show it to anyone.

> *Synthetic finding and data (safe to paste):* "In a (made-up) study, average reading time for a one-page brief fell as font size rose, then rose again when the font got too large. Synthetic data, mean seconds: 10 pt → 95 s; 12 pt → 78 s; 14 pt → 70 s; 16 pt → 72 s; 18 pt → 81 s."

First move. Ask a tool with a preview (for example, Claude artifacts or ChatGPT canvas) to build a small, self-contained HTML page with a labelled bar or line chart of these values, plus a one-sentence plain-language takeaway. Keep a **verification log**.

What to watch for. Do the **plotted numbers match** the data exactly? Is the **takeaway honest** about the U-shape (best around 14 pt), or does it over-simplify to "bigger is better"? Check **colour contrast** and that axes are labelled. Try a silly input (a negative font size) and see what breaks. "It rendered" is not "it is correct" — log every gap.

Track D · Public Engagement Translator

Scenario. Translate a hedged, cautious finding for three audiences — then audit each version for accuracy drift.

> *Synthetic abstract (safe to paste):* "In a small, preliminary, correlational study of 60 undergraduates, students who reported using a structured note-taking app **tended** to report slightly higher exam confidence (not exam *marks*). The effect was modest, the sample was not representative, and no causal claim can be made. Replication is needed."

First move. Produce three versions — an interested public, a policy reader and secondary-school pupils — then line each up against the source and mark every **over-claim, dropped caveat or false certainty**.

What to watch for. Watch the tool quietly upgrade "tended to report slightly higher **confidence**" into "**boosts exam results**", drop "small, preliminary, correlational", and add a confident headline the evidence cannot bear. Your **fidelity checklist** — "every claim traceable to the source; all caveats survive; no causal language" — is the artefact.

Starter-prompt library

Starting points, not scripts. Each is written to **build friction in** — to make the tool argue, flag uncertainty and expose its own weak spots, rather than produce a confident wall of text. Adapt freely; replace the bracketed parts; and remember the data decision aid before you paste anything.

Prompting for friction — five habits

1. **Make it disagree with itself.** Ask for the strongest case *against* its own answer.
2. **Force uncertainty to the surface.** "Mark each claim as confident, uncertain or a guess, and say what would change your mind."
3. **Demand checkable sources.** "Cite sources I can verify; if you are not sure a source exists, say so rather than inventing one."
4. **Ask what it is missing.** "What would an expert say you have overlooked?"
5. **Keep the human in the loop.** "List what a human must check before this is used."

Track A · Methodological Blind-Spot Detector

> "Act as a sceptical peer reviewer. Here is a study design: [PASTE ANONYMISED DESIGN]. Review it across these lenses — sampling and recruitment, common-method bias, construct validity, causal inference, generalisability and ethics. For each issue, rate how serious it is and explain *why*. Separate genuine, study-specific weaknesses from objections that would apply to almost any study. Do not invent citations."

> "Now argue the opposite: what is *methodologically strong* here, and which of your own criticisms is weakest?"

Track B · Accessible Executive-Function Layer

> "Turn these meeting notes into a structured action list with columns for action, owner, deadline and priority: [PASTE REDACTED NOTES]. Only include owners and dates that are explicitly stated. Put anything vague in a separate 'needs an owner / unclear' list. Do not invent details."

> "Review your own output: did you add any action, owner or date that was not in my notes? List exactly what you inferred versus what I gave you."

Track C · Rapid Prototyping for Knowledge Translation

> “Build a single self-contained HTML page that shows this data as a labelled chart, with a one-sentence plain-language takeaway: [PASTE SYNTHETIC DATA]. Use accessible colours and label the axes. Add a short comment noting any assumption you made.”

> “Now act as a reviewer of your own code: where could the numbers be wrong, what breaks with unexpected input, and what must a human verify before this is published?”

Track D · Public Engagement Translator

> “Rewrite this abstract for three audiences — an interested general reader, a policy reader and secondary-school pupils: [PASTE PUBLIC ABSTRACT]. Preserve every caveat and uncertainty. Do not add causal claims the abstract does not make.”

> “Audit your three versions against the original. List every place you strengthened a claim, dropped a caveat or added certainty the source did not have.”

General critical-use prompts (any track)

> “Before answering, tell me what you would need to know to answer well, and what assumptions you are making.”

> “Give your answer, then a section headed ‘Where I might be wrong’ and a section headed ‘What a human should check’.”

> “I think your answer is wrong because [REASON]. Defend it or revise it, and tell me which.”

> “Map the task: list its phases, mark what you would automate and where a human must steer or verify, then say whether oversight should be continuous (interwoven) or at checkpoints (staged) — and why.”

Disclaimer: these materials and the tools they reference are the facilitator's own choices, made in a personal capacity, and do not represent the views or official position of the University of Oxford, the facilitator's employer, or of the host, the University of Westminster. The University of Oxford accepts no liability for the selection, use or outcomes of any third-party tool. Participants remain solely responsible for compliance with their own institutional policy, the UK GDPR and research ethics. This wording is a template and is not legal advice.