

Group table pack

Group: _____ (a short, friendly name that does not identify anyone, such as 'Otters' or 'Team Kelp')

· **Track:** _____ (A · B · C · D)

Everything your group needs for the applied session is in this booklet. You have about fifteen minutes in Part 2. Get the core three done first. If that is all you manage, you have done the task.

1. One artefact. Run your tool on your problem and make something real: a critique, a template, a chart, a translation.
2. One or two caught errors. Moments it was fluent and wrong. Keep them.
3. One insight. The limitation, the safeguard or one honest thing this exposed about how your field is over- or under-using AI, which you will give in the lightning round, in about 45 seconds.

If you have about five minutes left, go deeper:

4. The automation–steering map: which steps the tool ran, and where you steered.
5. Interwoven or staged? Was your oversight woven through every step, or staged at checkpoints between phases? Say why, and what it cost.

A rough split that keeps you on time: about eight minutes to build, up to the pivot at about 13:26, then four to break it and name the insight, and the last three for the map and the oversight model if time is left. When you fall behind, drop the map before the insight.

The full rubric scoring and the societal reflection are in this pack for your reflection and the lightning round, and they are not part of the fifteen minutes. A group that wants to record them can write a HackMD or paper note, which the facilitator archives alongside the app export.

Before you paste anything: the red lines

Never put personal data, students' marks, references or pastoral notes, unpublished or participant data, or confidential, peer-review or embargoed material into a free AI tool. When in doubt, anonymise, synthesise or abstain. The one-page data decision aid in this pack walks you through it.

Your shared note

Your group's note lives in the workshop app. Scan the QR code or go to genai-rt.web.app, and your facilitator reads out the session passcode. One device starts the group, and the others join with the code it shows. Fill in the essentials as you go, and submit before the lightning round at 13:33 where you can, and by 13:45 at the latest. There is no account and nothing to install.

One note, one surface. Use the app by default. If you have no device or it will not connect, HackMD (hackmd.io) is the fallback, with the same headings. Do not use both.

One reflection for the wrap-up

This is for the tidy-up or the lightning round, and not for the fifteen minutes. Use it as your lightning-round line if it is the stronger one: is your field generally under- or over-using generative AI for this kind of task? Why, with what consequences, and how would you redress it over the next two years?

What is in this pack

- Role cards: cut along the lines and take one each (five angles, one per person).
 - Data decision aid: the 'can I paste this?' one-pager.
 - The rubric: what you are working towards, and the fuller note for the fallback.
 - A worked example per track: your default problem if no one brought a real one.
 - Starter prompts: copy, adapt and go.
-

A personal selection, made in the facilitator's own capacity. It is not the position of the University of Oxford (the facilitator's employer) or of the host, and it is not legal advice.

Group role cards (optional prompt sheet)

Five angles to keep an eye on in a group of five. Take one each, or divide them however suits you. They are about where attention goes, and never about who owns the work, so rotate the keyboard freely. Each maps to one rubric dimension, so make sure none is left unwatched. An experienced group can ignore these entirely, provided that the building does not crowd out the interrogating.

Print tip. Optional. Print one set per group if it helps, and cut along the rules. At four people, the Convenor doubles as Steward. At six, add a second Sceptic, since more red-teaming does no harm.

The mapping (canonical):

Angle (handle)	Rubric dimension	In one line
Convenor	Project Definition	Keeps the real problem and the clock in focus
Driver	Technology Stack	Operates the tools and owns the tool choices
Steward	Data Security & Ethics	Guards the red lines
Reporter	Financial & Scalability	Owns the record, the submission and the 45-second insight
Sceptic	Human-in-the-Loop Protocol	Tries to break the output and curates the caught errors



CONVENOR · *Project Definition*

Your angle: the group stays on a real problem and on time.

- **Do:** in the first minute, agree your one problem (the seed you formed around, or a real one a member brought) and a success criterion; split the five angles; watch the clock against the timeline; make sure everyone speaks.
- **At the pivot (about 13:26):** when the facilitator calls the halt, take your group through the four questions on screen and move it from building to interrogating.
- **Ask often:** 'What is the actual problem, and how will we know it helped?'



DRIVER · *Technology Stack*

Your angle: operating the tools and making the technical choices.

- **Do:** share your screen; place your tool on the spectrum (off-the-shelf, no-code or IDE-API); be ready to say why this level and not one up or down.

- **At the pivot (about 13:26):** note what the tool made easier, and the cost.
- **Ask often:** ‘Would a different level of tool change what we can trust?’



STEWARD · *Data Security & Ethics*

Your angle: holding the red lines.

- **Do:** vet everything before it is pasted (use the data decision aid); watch for personal, special-category or confidential data; track disclosure, IP and consent.
- **At the pivot (about 13:26):** confirm nothing crossed a red line, and flag anything to redact.
- **Ask often:** ‘Whose data is this, and would we be content for it to be public?’



REPORTER · *Financial & Scalability*

Your angle: the record and the insight.

- **Do:** keep the group's note in the app (or in the rubric template on the fallback); judge whether the approach survives the free tier and scale; submit the note for review before the lightning round at 13:33 where you can, and by 13:45 at the latest; give the 45-second insight at 13:33 (no slides).
- **At the pivot (about 13:26):** make sure the note is being filled in as you go, and not left to the end.
- **Ask often:** ‘Does this still work next month, for fifty people, for free?’



SCEPTIC · *Human-in-the-Loop Protocol*

Your angle: breaking the output, and deciding where the human must stay.

- **Do:** stress-test every result; hunt hallucinations, bias and false confidence; curate the museum of caught errors; name the checkpoints and what you will never delegate; draw the automation–steering map and decide whether oversight is interwoven or staged.
- **At the pivot (about 13:26):** bring your best caught error to the group.
- **Ask often:** ‘Where, exactly, does a human have to check before we trust this?’



A personal selection, made in the facilitator's own capacity. It is not the position of the University of Oxford (the facilitator's employer) or of the host, and it is not legal advice.

One-page data decision aid: can I paste this in?

Print one per group, or keep it open in a tab, and use it before anything goes into a free, third-party AI tool. It is a thinking aid and not legal advice. Institutional policy, the UK GDPR and research ethics still govern.

The one-line test

If you would not pin it to a public noticeboard in the building, do not paste it into a free consumer AI tool.

Free tiers commonly reuse your input to train models, and settings change without notice. Treat anything you type as potentially non-private and non-retractable.

Walk the four questions

1. Is it personal data? Does it relate to an identifiable living person (names, emails, student or participant records, anything that could single someone out, even indirectly)? If yes, do not paste it as it is. Anonymise or synthesise it, and it is then no longer personal data. If it has to stay identifiable, a lawful basis is necessary but nowhere near sufficient: you also need a tool your institution has cleared, with a processor contract behind it, which a free consumer tool is not.
2. Is it special category data under Article 9 of the UK GDPR? That means racial or ethnic origin, political opinions, religious or philosophical beliefs, trade union membership, genetic data, biometric data used to identify someone, health, sex life or sexual orientation. Criminal offence data, under Article 10, carries similar weight. If yes, stop. A free consumer tool is the wrong place. Anonymise beyond recognition, synthesise or abstain. Children's data and commercially sensitive material are not special category, but both carry extra duties, so treat them the same way here.
3. Is it confidential, unpublished or restricted? This covers unpublished data or manuscripts, grant drafts, peer-review files, material covered by a non-disclosure agreement, anything under embargo and anything shared without co-author or participant consent. If yes, stop. Confidentiality and embargoes are not yours alone to waive.
4. Do you hold the rights? Is the text or image yours to share, or is it someone else's intellectual property? If unsure, treat it as restricted.

If you cleared all four, you are probably fine. If any gave you pause, that pause is the signal, so act on it.

Green, amber and red at a glance

● Generally fine	● Pause and treat with care	● Do not paste into a free tool
Already-public text (your published abstract)	Lightly disguised real material	Personal data of identifiable people
Fully synthetic or fabricated examples	Aggregated or partially de-identified data	Special category, criminal offence or children's data
Generic, non-identifying questions	Draft prose with no sensitive content	Unpublished data, grants, peer review
Your own writing, where you hold the rights	Anything you would hesitate to email widely	Confidential, embargoed or NDA-covered material

If the answer is no: what to do instead

- Anonymise. Remove names, identifiers and anything that singles someone out, and check that it cannot be re-identified by combination.
- Synthesise. Recreate the structure of the problem with invented content. A critique of a method works just as well on a paraphrased method.
- Abstract up a level. Discuss the shape of the task, and leave out the sensitive specifics.
- Use an appropriately governed tool. An enterprise or paid service under a proper data-processing agreement may be acceptable where a free consumer tool is not. That is a decision for your institution, and it is out of scope today.
- Abstain. Sometimes the right answer is to leave the tool alone. That is a valid finding, and a good one for your museum of caught errors.

Disclaimer: these materials and the tools they reference are the facilitator's own choices, made in a personal capacity, and do not represent the views or official position of the University of Oxford, the facilitator's employer, or of the host, the University of Westminster. The University of Oxford accepts no liability for the selection, use or outcomes of any third-party tool. Participants remain solely responsible for compliance with their own institutional policy, the UK GDPR and research ethics. This wording is a template and is not legal advice.

Evaluation rubric and note template

This is the group's workbench for the day. By default you capture your work in the workshop app (genai-rt.web.app), whose short form covers the essentials: the problem, the artefact, the errors you caught, the automation–steering map, the oversight model, your insight and a field reflection. Nothing needs setting up beyond the passcode. The fuller template below adds the five rubric dimensions, each scored from 1 to 5 (1 Nascent, 3 Developing, 5 Robust), a deliverable link, the museum table and the societal reflection. It is the note for the HackMD or paper fallback, and for a group that wants to record its scores and reflection after submitting in the app.

You are not graded. The rubric is a thinking tool and a shared record, and, once archived, part of the open account of what this cohort learned.

A note on the short format. The build window in Part 2 is about 15 minutes. Get the core three first: the artefact (your tool run on your problem), one caught error and one insight for the lightning round. If you have about five minutes left, add the automation–steering map and say whether your oversight was interwoven or staged. The rubric scores and the societal reflection come later, when you tidy the note. A thoughtful partial note is worth more than a rushed full one.

A. The HackMD fallback (about five minutes, Reporter drives, only if you cannot use the app)

1. Go to hackmd.io and sign in. It is free, and you can use a Google, GitHub or email login. There is nothing to install.
2. Click '+ New note'. The left pane is Markdown and the right pane shows a live preview. You can ignore the formatting and type straight into the headings.
3. Copy the whole template in section B below (everything inside the box) and paste it into the note.
4. Fill it in as you work, and avoid leaving it all to the end. Rough notes are fine.
5. Make it link-shareable: open the note's sharing or permissions control and set it so that anyone with the link can read (you do not need to grant editing).
6. Copy the note's link.
7. Share that link with the facilitator before the lightning round at 13:33 where you can, and by 13:45 at the latest. Show the note on screen with your group number and track. You do not need a GitHub account, because the facilitator gathers every group's link.

If HackMD will not let you sign up (some logins are restricted), do not lose time troubleshooting: switch to paper at once. Write the same template in any shared document or on paper. It is plain Markdown, and the facilitator can transcribe one note into the archive afterwards.

After the session, the facilitator exports every note whose group opted in to sharing, removes anything that should not be public, and commits it under `submissions/`.

B. The note template (paste this box into HackMD or copy it on paper; the

app has its own shorter form)

[Group NN] · Track [A / B / C / D] · 24 June 2026

Roles taken (no names needed): Convenor · Reporter · Driver · Sceptic · Steward

The real problem we brought: ...

Share publicly? yes / no (yes means the group is content for this note, which should name no one, to enter the public archive)

1. Project Definition

- The real problem, in one sentence a colleague would recognise: ...
- Success criterion (how we would know the tool actually helped): ...

2. Technology Stack

- Tool(s) used, and where they sit on the spectrum (off-the-shelf / no-code / IDE-API): ...
- Why this level and not one up or down (control, data exposure, effort): ...

3. Data Security & Ethics

- What we put into the tool, and how we de-identified it: ...
- Red lines we held; UK GDPR, ethics, disclosure of AI use, IP, consent and fairness: ...

4. Financial & Scalability Constraints

- Free-tier limits we hit or foresee: ...
- What breaks first at scale (cost, usage limits or trust): ...

5. Human-in-the-Loop Protocol

- Automation-steering map. Break the project into its phases; for each, note what is automated (the tool does it) and what is human-steered (we decide, verify or override):

Phase	Automated (AI does)	Human steering (we decide / verify / override)
---	---	---
...	...	...

- Interwoven or staged? Is our oversight interwoven (continuous, with a human in the loop at every step) or staged (concentrated at distinct checkpoints between phases)? Which did we choose, and why?
- What our choice costs, and how we offset it. (Interwoven oversight is thorough but heavy, and checking every step can dull attention until things are waved through out of habit; staged oversight is lighter, but errors can accumulate unnoticed between checkpoints.): ...
- Checkpoints we set in advance: ...
- What we will never delegate to the tool: ...
- Who is accountable for the final judgement: ...

Deliverable

- Link or screenshot of the artefact we made: ...

Museum of caught errors

What happened	How we caught it	What it signals about where humans must stay
---	---	---
...	...	...

Societal reflection (complete when you tidy and share the note; no score, just think)

- Thinking versus writing: if the tool did the writing, what thinking did we still have to do,

and what might we lose by offloading it? ...

- Fairness: could this widen or narrow disparities (first versus second language, more versus less support, career stage, socioeconomic background)? Who gains, and who is left behind? ...
- Disclosure: would we disclose using AI here? Would we, or should we, disclose the conditions it offset (for example, writing in a second language or limited support)? Why or why not? ...
- Disciplinary norms: is our field generally under- or over-using generative AI for this kind of task? Why, with what consequences, and how would we redress it over the next two years? ...

Lightning-round insight (45 seconds, no slides)

- > Pick the strongest of three: the single most significant limitation we found; the most
- > important human-in-the-loop safeguard we built in; or one honest thing this exposed
- > about how our field is over- or under-using AI for this task:
- >
- > ...

Self-assessment (complete when you tidy and share the note; 1 Nascent · 3 Developing · 5 Robust)

Dimension	Score (1–5)	One-line justification
Project Definition		
Technology Stack		
Data Security & Ethics		
Financial & Scalability		
Human-in-the-Loop		

C. The rubric anchors

Use these to place yourselves honestly. Most groups will sit around 3 on most dimensions in a short applied session, and that is as it should be. A thoughtful 3 with a sharp caught error is worth more than an unexamined 5. Use 2 and 4 for work that sits between two anchors: 2 is nearer Nascent than Developing, and 4 is nearer Robust than Developing.

1 · Project Definition

Is the real problem, and its success criterion, clear?

- **1, Nascent:** the task is vague or a toy ('use AI on something'), with no success criterion.
- **3, Developing:** a real problem is named and there is a rough sense of what 'better' would look like, although the scope may be broad.
- **5, Robust:** a specific, bounded, genuinely held problem a colleague would recognise, with an explicit success criterion you could actually test.

2 · Technology Stack

Is it the right tool on the spectrum, and why?

- **1, Nascent:** one tool used by default, with no awareness of alternatives or of the spectrum.
- **3, Developing:** the choice is explained, with some sense of where it sits and what moving a level up or down would change.

- **5, Robust:** deliberate placement on the spectrum, with a reasoned trade-off of control, data exposure and effort, and alternatives weighed.

3 · Data Security & Ethics

Were the red lines held, and were the UK GDPR, ethics and fairness addressed?

- **1, Nascent:** personal, special-category or confidential data used without thought, and no consideration of disclosure, consent or fairness.
- **3, Developing:** red lines mostly respected and data de-identified, with some ethical reflection on disclosure, IP, bias or fairness.
- **5, Robust:** red lines explicitly held; anonymisation or lawful basis reasoned; disclosure, IP, consent, bias and fairness addressed; and a note that would withstand being public.

4 · Financial & Scalability Constraints

Does the free tier hold, and what breaks at scale?

- **1, Nascent:** no thought to cost or scale, and an assumption that it stays free forever.
- **3, Developing:** free-tier limits noted, with a rough sense of what scaling would cost or require.
- **5, Robust:** clear about free-tier viability, lock-in and the precise point at which data, cost or trust breaks the approach.

5 · Human-in-the-Loop Protocol

Are there explicit checkpoints, verification and accountability, and a deliberate map of automation against human steering?

- **1, Nascent:** output trusted as it is, with no checkpoint, no named accountability and no sense of where automation ends and human steering begins.
- **3, Developing:** some verification done, one or two checkpoints identified and an implicit split between the automated and human-steered steps.
- **5, Robust:** a deliberate automation–steering map; oversight consciously chosen as interwoven or staged, with the weakness of that choice offset; an explicit list of what is never delegated; and a named person who owns the final judgement on correctness, ethics and attribution.

Disclaimer

These materials and the tools they reference are the facilitator's own choices, made in a personal capacity. They do not represent the views or official position of the University of Oxford, the facilitator's employer, or of the host, the University of Westminster. The University of Oxford accepts no liability for the selection, use or outcomes of any third-party tool. Participants remain solely responsible for compliance with their own institutional policy, the UK GDPR and research ethics.

This wording is a template and is not legal advice.

Worked examples: one per track

A worked example is the fallback, and not the default. Use your group's seed, or a real, non-confidential problem a member brings, and turn to the example below only if neither is to hand. It is also what you use when the Wi-Fi fails and you switch to critiquing on paper.

The material below is fully synthetic and safe to paste, with no real people, data or unpublished work, so a group with nothing to hand can start at once. If you are using it, still spend a minute on the data decision aid: say what you would have had to strip out had this been your own material, and record that under Data Security & Ethics. Each example is sized for the short build window and leads naturally to a caught error.

Track A · Methodological Blind-Spot Detector

Scenario. You are about to submit the study below. Use an AI tool as a critical reviewer, then verify which of its objections are real, which are generic and which are wrong.

Synthetic methods snippet (safe to paste): 'We ran a cross-sectional online survey of 200 marketing managers recruited through a LinkedIn post. Each manager completed a 12-item self-report scale of "authentic leadership" and, in the same questionnaire, rated their own team's "innovativeness" on a 5-item scale. Authentic leadership correlated positively with team innovativeness ($r = .46$, $p < .001$). We conclude that authentic leadership causes higher team innovation and recommend leadership training to raise innovation.'

First move. Ask the tool to review across explicit lenses: sampling and recruitment, common-method bias, construct validity, causal inference and generalisability. Then classify each point it raises as real, generic or wrong.

What to watch for. A strong tool should flag the causal claim from cross-sectional data, common-method or single-source bias (the same person rates both variables), self-selection via LinkedIn and self-report of one's own team. Note whether it also pads the list with vague, always-true objections, and whether it ever invents a statistic or a citation. Those are your museum pieces.

Track B · Accessible Executive-Function Layer

Scenario. Turn a messy set of meeting notes into a clean, owned action list, then check that the tool did not quietly invent or misattribute anything.

Synthetic meeting notes (safe to paste; roles only, no real people): 'Centre meeting. Talked about the seminar series. Someone should sort speakers for the autumn, maybe the deputy director? Budget underspend, need to use it before July or we lose it. The new PhD reps raised workload worries. Website is out of date (old staff list). Agreed to revisit the mentoring scheme but no one said who. Ethics turnaround is too slow. Chase the committee. Next meeting in three weeks, same room hopefully.'

First move. Ask for a structured action list with owner, deadline and priority, plus a separate 'unclear or needs an owner' list. Build it as a reusable template, and not as a one-off.

What to watch for. Does the tool invent owners or dates that the notes never stated ('Deputy Director, by 30 June')? Does it silently drop the items with no owner instead of flagging them? Does it confidently over-formalise a vague discussion? The safeguard you design ('every action must trace to a line in the notes; no invented owners') is the deliverable.

Track C · Rapid Prototyping for Knowledge Translation (*technical stretch*)

Scenario. Build a tiny interactive that communicates a finding, and verify that it is actually correct before you would ever show it to anyone.

Synthetic finding and data (safe to paste): 'In a (made-up) study, average reading time for a one-page brief fell as font size rose, then rose again when the font got too large. Synthetic data, mean seconds: 10 pt → 95 s; 12 pt → 78 s; 14 pt → 70 s; 16 pt → 72 s; 18 pt → 81 s.'

First move. Ask a tool with a preview (for example, Claude artifacts or Gemini's Canvas) to build a small, self-contained HTML page with a labelled bar or line chart of these values, plus a one-sentence plain-language takeaway. Keep a verification log.

What to watch for. Do the plotted numbers match the data exactly? Is the takeaway honest about the U-shape (fastest reading around 14 pt), or does it over-simplify to 'bigger is better'? Check colour contrast and that the axes are labelled. Then add a sixth point of your own, 20 pt, and see whether the axis, the scale and the takeaway update correctly or whether the tool hard-coded them. The fact that it rendered does not mean that it is correct, so log every gap.

Track D · Public Engagement Translator

Scenario. Translate a hedged, cautious finding for three audiences, then audit each version for accuracy drift.

Synthetic abstract (safe to paste): 'In a small, preliminary, correlational study of 60 undergraduates, students who reported using a structured note-taking app tended to report slightly higher exam confidence (not exam marks). The effect was modest, the sample was not representative, and no causal claim can be made. Replication is needed.'

First move. Produce three versions (an interested public, a policy reader and secondary-school pupils), then line each up against the source and mark every over-claim, dropped caveat or false certainty.

What to watch for. Watch the tool quietly upgrade 'tended to report slightly higher confidence' into 'boosts exam results', drop 'small, preliminary, correlational' and add a confident headline the evidence cannot bear. Your fidelity checklist ('every claim traceable to the source; all caveats survive; no causal language') is the artefact.

Disclaimer: these materials and the tools they reference are the facilitator's own choices, made in a personal capacity, and do not represent the views or official position of the University of Oxford, the facilitator's employer, or of the host, the University of Westminster. The University of Oxford accepts no liability for the selection, use or outcomes of any third-party tool. Participants remain solely responsible for compliance with their own institutional policy, the UK GDPR and research ethics. This wording is a template and is not legal advice.

Starter-prompt library

Starting points, and never scripts. Adapt freely, replace the bracketed parts and remember the [data decision aid](#) before you paste anything.

Each track has two prompts, and the order matters. Run the first pass as it stands, without the guardrails, because a prompt that forbids a failure mode also hides it, and the failure is what you are here to catch. Then design the second prompt as the safeguard your group would actually adopt, and note the difference between them. That difference is your finding.

Prompting for friction: five habits

1. Make it disagree with itself. Ask for the strongest case against its own answer.
 2. Ask it to sort its own claims. 'Mark each claim as confident, uncertain or a guess, and say what would change your mind.' These labels are the model's guess about itself and not a measurement, and a confidently wrong claim will be labelled confident, so use them only to decide what to verify first, and verify it elsewhere.
 3. Demand checkable sources. 'Cite sources I can verify. If you are not sure a source exists, say so instead of inventing one.'
 4. Ask what it is missing. 'What would an expert say you have overlooked?'
 5. Keep the human in the loop. 'List what a human must check before this is used.'
-

Track A · Methodological Blind-Spot Detector

First pass, unguarded. Naming the flaws for it would tell you only that it can follow a list, so do not.

'Act as a sceptical peer reviewer. Here is a study design: [PASTE ANONYMISED DESIGN]. What is wrong with it? For each issue, say how serious it is and why.'

Then classify every objection yourselves as real, generic or wrong, and check any factual claim outside the tool. Only now add the guardrails, as the protocol you would give a colleague.

'Review it again across these lenses: sampling and recruitment, common-method bias, construct validity, causal inference, generalisability and ethics. Separate study-specific weaknesses from objections that would apply to almost any study, and do not invent citations.'

'Now argue the opposite: what is methodologically strong here, and which of your own criticisms is weakest?'

Track B · Accessible Executive-Function Layer

First pass, unguarded. Watch what it does with the items that have no owner and no date.

‘Turn these meeting notes into a structured action list with columns for action, owner, deadline and priority: [PASTE REDACTED NOTES].’

Check every owner and date against the notes. Then write the safeguard, which is the deliverable for this track.

‘Do it again. Only include owners and dates that are explicitly stated, put anything vague in a separate list headed “needs an owner or unclear”, and do not invent details.’

‘Review your own output: did you add any action, owner or date that was not in my notes? List exactly what you inferred and what I gave you.’

Track C · Rapid Prototyping for Knowledge Translation

‘Build a single self-contained HTML page that shows this data as a labelled chart, with a one-sentence plain-language takeaway: [PASTE SYNTHETIC DATA]. Use accessible colours and label the axes. Add a short comment noting any assumption you made.’

‘Now act as a reviewer of your own code: where could the numbers be wrong, what breaks with unexpected input, and what must a human verify before this is published?’

Track D · Public Engagement Translator

First pass, unguarded. Telling it to keep the caveats would hide the drift you are looking for.

‘Rewrite this abstract for three audiences, an interested general reader, a policy reader and secondary-school pupils: [PASTE PUBLIC ABSTRACT].’

Line each version up against the source yourselves and mark every over-claim, dropped caveat and added certainty. Then write the checklist as a prompt.

‘Do it again, preserving every caveat and uncertainty, and adding no causal claim the abstract does not make.’

‘Audit your three versions against the original. List every place you strengthened a claim, dropped a caveat or added certainty the source did not have.’

General critical-use prompts (any track)

‘Before answering, tell me what you would need to know to answer well, and what assumptions you are making.’

‘Give your answer, then a section headed “Where I might be wrong” and a section headed “What a human should check”.’

‘I think your answer is wrong because [REASON]. Defend it or revise it, and tell me which.’

‘Map the task: list its phases, mark what you would automate and where a human must steer or verify, then say whether oversight should be continuous (interwoven) or at checkpoints (staged), and why.’

Disclaimer: these materials and the tools they reference are the facilitator’s own choices, made in a personal capacity, and do not represent the views or official position of the University of Oxford, the facilitator’s employer, or of the host, the University of Westminster. The University of Oxford accepts no liability for the selection, use or outcomes of any third-party tool. Participants remain solely responsible for compliance with their own institutional policy, the UK GDPR and research ethics. This wording is a template and is not legal advice.