

Generative AI in research and teaching

Workshop print bundle

Item	Copies	Sides	Paper
Part 1 — cut these			
Role cards, for cutting (3 pages)	10 copies	SINGLE-SIDED	A4 portrait
Facilitator cue cards, for cutting (8 pages)	1 copy	SINGLE-SIDED	A4 portrait
Part 2 — collate into group bundles			
Group pack booklet (23 pages)	10 copies, plus one spare	Double-sided	A4 portrait
Group quick-start one-pager (1 page)	20 copies	Single-sided	A4 portrait
Part 3 — facilitator only			
Run sheet and morning checklist (6 pages)	1 copy	Double-sided	A4 portrait
Lightning-round tally (1 page)	1 copy	Single-sided	A4 portrait
Day-of app reset (2 pages)	1 copy	Single-sided	A4 portrait
Part 4 — put up around the room			
Seed-idea signs (10 pages)	1 copy	SINGLE-SIDED	A4 LANDSCAPE
Table numbers (10 pages)	1 copy	SINGLE-SIDED	A4 LANDSCAPE

This cannot be run as one job with one setting. The copy counts differ per item, the two card sections must print single-sided because they are cut, and the room signs are landscape. Each section below opens with its own instruction slip; the slips are not part of the job and need not be printed.

Part 1 — cut these

Role cards, for cutting

Copies	10 copies
Sides	SINGLE-SIDED
Paper	A4 portrait
Finishing	Cut once across the middle of every sheet
Length	3 pages

5 cards, two to a sheet. Every sheet cuts at the same height, so the whole stack can go through a guillotine in one pass. Cutting a double-sided sheet would destroy its reverse, which is why this section prints on one side only.

Convenor · *Project Definition*

Your angle is that the group stays on a real problem, and on time.

- **In the fifteen minutes.** Settle on your one problem in the first minute, either the seed your group formed around or a real one someone brought, and say how you would know the tool had helped. Share the five angles out, keep half an eye on the clock and make room for everyone to speak.
- **At the pivot, about 13:26.** The halt lasts about half a minute, so keep it short. Put the four questions on screen to the group and move it from building to examining.
- **Worth asking.** 'What is the actual problem here, and how will we know it helped?'

Role card 1 of 5

Driver · *Technology Stack*

Your angle is running the tool, and the choices that go with it.

- **In the fifteen minutes.** Share your screen so the group can see what the tool is doing. Place your tool on the spectrum, whether off-the-shelf, no-code or at the IDE and API level, and be ready to say why you chose that level over one above or below.
- **At the pivot, about 13:26.** Have an answer ready, in a sentence, for what the tool genuinely made easier and what that ease cost.
- **Worth asking.** 'Would a different level of tool change what we can trust?'

Role card 2 of 5

Steward · *Data Security and Ethics*

Your angle is holding the red lines.

- **In the fifteen minutes.** Look over anything before it is pasted, using the data decision aid in your pack. Watch for personal data, special-category data and confidential material, and keep track of disclosure, intellectual property and consent as they come up.
- **At the pivot, about 13:26.** Satisfy yourself that nothing has crossed a red line, and speak up only if something needs redacting.
- **Worth asking.** 'Whose data is this, and would we be content for it to be public?'

Role card 3 of 5

Reporter · *Financial and Scalability*

Your angle is the record, and the insight the room hears.

- **In the fifteen minutes.** Keep the group's note in the app as you go, or in the rubric template if you are on the paper fallback. Form a view on whether the approach survives the free tier and a larger cohort. Submit the note before the lightning round at 13:33 where you can, and during it at the latest.
- **At the pivot, about 13:26.** Glance at the note and check it is being filled in as you go, instead of leaving it to the last minute.
- **Worth asking.** 'Does this still work next month, for fifty people, for free?'

Role card 4 of 5

Sceptic · *Human-in-the-Loop Protocol*

Your angle is testing the output, and deciding where the human has to stay.

- **In the fifteen minutes.** Press on every result, looking for invented detail, bias and confident error, and keep the good catches for the museum. Name the checkpoints the group would want and the things it would never delegate. If there is time, sketch the automation and steering map, and decide whether your oversight ran throughout or at set points.
- **At the pivot, about 13:26.** This is the one moment to speak: bring your best caught error to the group.
- **Worth asking.** 'Where, exactly, does a human have to check before we trust this?'

Role card 5 of 5

Blank. Keep it for notes.

Facilitator cue cards, for cutting

Copies	1 copy
Sides	SINGLE-SIDED
Paper	A4 portrait
Finishing	Cut once across the middle of every sheet, then keep in clock order
Length	8 pages

15 cards, two to a sheet, running from 12:00 to 13:40. Each card is numbered in its corner, so the order survives the cutting.

12:00 · WELCOME AND STANCE (Part 1)

- **Say:** 'We are here for discernment, not enthusiasm: a deeper insight into the technology, and no training on any one tool. Whatever your confidence, you are welcome.'
- **Do:** read the disclaimer aloud, and make the point that friction is a signal.
- **Watch:** start on time. This is the conceptual part, so keep it brisk.

Cue card 1 of 15

12:03 · SAVED OR BURNED

- **Say:** 'One sentence each with your neighbour: the last time an AI tool saved you or burned you.'
- **Do:** two minutes, in parallel, then pull the room back.

Cue card 2 of 15

12:05 · A LONG LINE OF TOOLS

- **Say:** 'AI is the latest in a long line of productivity tools, from the calculator and the computer to the spreadsheet and now generative AI, in an economy that rewards output. The question is how to use it with discernment.'
- **Do:** show the chain on screen, and take no more than a couple of minutes.

Cue card 3 of 15

12:08 · GETTING STARTED (no experience needed)

- **Say:** 'A chat assistant: you type a request and it writes back. Pick one free tool. Say who you are, what you want and what good looks like, then push back.'
- **Do:** point to the starter prompts, and reassure the room that mixed groups are a strength and no one is behind.

Cue card 4 of 15

12:10 · FRICTION IS A SIGNAL

- **Say:** 'When the tool resists you, that friction is information. It marks where your judgement does the work.'
- **Do:** Lee et al. and Drosos et al. in a sentence each. Do not lecture.

Cue card 5 of 15

12:14–12:21 · SPECTRUM · RED LINES · ETHICS

- **Do:** one breath each (slides 10–12). Name the teaching angle: students' data, assessment integrity, disclosure.
- **Watch:** these are reference frames for Part 2, and there is no need to master them now.

Cue card 6 of 15

12:21 · THE BIGGER QUESTIONS AND THE HUMAN IN THE LOOP

- **Say:** 'Zoom out: the shift between thinking and writing; fairness, meaning who is helped and who is left out by language, support, career stage or background; and disclosure. There are no answers today, so notice them in Part 2.'
- **Do:** then the human in the loop and the map (slide 14). Accountability stays human.

Cue card 7 of 15

12:26 · WHAT PART 2 ASKS, THEN FORM GROUPS (end of Part 1)

- **Say:** 'These ideas, and the bigger questions, are the rubric and the reflection you will fill in. Now form groups of five: walk to the sign for the idea you would most like to try, and the letter is your track.'
- **Do:** preview the rubric (slide 15) and the tracks (slide 16), then send them to lunch. End Part 1 at 12:30.

Cue card 8 of 15

12:30–13:15 · NETWORKING LUNCH (not your session)

- **Do:** a networking break. Use part of it for the app reset if it is not already done. Part 2 starts at 13:15.

Cue card 9 of 15

13:15 · SETTLE IN (Part 2)

- **Say:** 'Re-find your group. Quickly agree your one problem, the seed or a real one someone brought. Open the app at genai-rt.web.app. Here is the passcode. Red lines on.'
- **Do:** about three minutes in all: regroup, agree the one problem in roughly a minute, then open the app and read out the passcode. Keep it tight so that it does not eat into the fifteen-minute build window.

Cue card 10 of 15

13:18 · APPLY IT (15 min)

- **Do:** let them work. Stop sensitive data, unstick stalls and offer one provocation each.
- **Push:** one caught error and one insight first. Where a group has time, the automation–steering map and the choice between interwoven and staged oversight.

Cue card 11 of 15

About 13:26 · PIVOT 🕒

- **Say:** ‘Thirty seconds, keyboards down. What did it make easier, and at what cost? Where did it resist you? Have you held the red lines? Now switch from building to interrogating.’
- **Say, while every form is still open:** ‘One more thing before you submit. If you are happy for your note to join the public archive, tick the sharing box at the foot of the form. It is optional, and unticked is a perfectly good answer. Submitting locks the box, so decide it now.’

Cue card 12 of 15

13:33 · LIGHTNING ROUND (7 min)

- **Say:** 'One insight each and no slides. I will call the time. Give us the biggest limitation, the best safeguard or how your field over- or under-uses AI here.'
- **Do:** visible timer. 45 seconds each with eight groups, or announce 35 seconds if there are nine or ten, since each handover costs five to ten seconds. Still behind after six? Take one line from the rest. Write one line per group on the tally sheet, and leave every approval until after the session.

Cue card 13 of 15

13:40 · CLOSE (5 min)

- **Do:** synthesise the round by naming two or three threads from it.
- **Connect:** friction as signal, the spectrum, the red lines, the human in the loop.
- **Give:** the take-home actions, the further reading and the follow-up.
Close at 13:45.

Cue card 14 of 15

Any time · IF IT BREAKS

- **Wi-Fi or the app is down.** The pack is enough. Groups critique the printed worked example against the rubric, on paper, and you type the notes up afterwards.
- **Projection fails.** Everything a group needs is on the one-pager already on its table.
- **Someone arrives for Part 2 only.** Put them into any group of four, hand them a one-pager and let the group brief them in a sentence. No catching up is needed.
- **Part 1 starts late.** Drop 'Saved or burned', then compress slides 10 to 13. Never cut the group formation at 12:26, since Part 2 depends on it.
- **A group is smaller than four or larger than six.** Below four, merge it with a neighbour. Above six, split it in two and give each half a lightning-round slot.

Cue card 15 of 15

Blank. Keep it for notes.

Part 2 — collate into group bundles

Group pack booklet

Copies	10 copies, plus one spare
Sides	Double-sided
Paper	A4 portrait
Finishing	Stapled, top left
Length	23 pages

One booklet per table. The cover has fill-in fields, so every table needs its own. The footer numbers every page, so a booklet missing a leaf shows up while collating.

Group table pack

Group: _____ (a short, friendly name that does not identify anyone, such as 'Otters' or 'Team Kelp') · **Track:** _____ (A · B · C · D)

Table we return to at 13:15: _____ · **Join code the app shows** (after 13:15): _____

Everything your group needs for the applied session is in this booklet. You have about fifteen minutes in Part 2. Get the core three done first. If that is all you manage, you have done the task.

1. One artefact. Run your tool on your problem and make something real: a critique, a template, a chart, a translation.
2. One or two caught errors. Note the moments when it was fluent and wrong, and keep them.
3. One insight. The limitation, the safeguard or one honest thing this exposed about how your field is over- or under-using AI, which you will give in the lightning round, in up to 45 seconds.

If you have about five minutes left, go deeper:

4. The automation—steering map: which steps the tool ran, and where you steered.
5. Interwoven or staged? Was your oversight woven through every step, or staged at checkpoints between phases? Say why, and what it cost.

A rough split keeps you on time. Give about eight minutes to building, up to the pivot at about 13:26, and the seven after it to testing what you made, naming your insight and submitting. Add the map and the oversight model in that stretch if five minutes or so remain. When you fall behind, drop the map and never the insight.

The app has a box for the field reflection, so record it there if you get to it. The five rubric scores are in this pack alone, and they belong to the tidy-up afterwards, once the fifteen minutes are over.

Before you paste anything: the red lines

Never put personal data, students' marks, references or pastoral notes, unpublished or participant data, or confidential, peer-review or embargoed

material into a free AI tool. When in doubt, anonymise, synthesise or abstain. The one-page data decision aid in this pack walks you through it.

Your shared note

Your group's note lives in the workshop app. Scan the QR code on screen or go to genai-rt.web.app, and your facilitator reads out the session passcode. One device starts the group, and the others join with the code it shows. Fill in the essentials as you go, and submit before the lightning round at 13:33 where you can, and during it at the latest. There is no account and nothing to install.

Keep to one note on one surface, and use the app by default. If nobody at the table has a device, work on paper: the worked example and the rubric in this pack are all you need, and your facilitator will type it up afterwards. If you have a device but it will not connect, HackMD (hackmd.io) is the fallback and carries the same headings, so please keep to whichever one you start with.

One reflection for the wrap-up

This belongs in the tidy-up or the lightning round rather than the fifteen minutes. Use it as your lightning-round line if it is the stronger one: is your field generally under- or over-using generative AI for this kind of task? Why, with what consequences and how would you redress it over the next two years?

What is in this pack

- Role cards: cut along the lines and take one each (five angles, one per person).
 - Data decision aid: the 'can I paste this?' one-pager.
 - The rubric: what you are working towards, and the fuller note for the fallback.
 - A worked example per track: your default problem if no one brought a real one.
 - Starter prompts: copy, adapt and go.
-

A personal selection, made in the facilitator's own capacity. It is not the position of the University of Oxford (the facilitator's employer) or of the host, and it is not

legal advice.

Group role cards (optional prompt sheet)

Five angles to keep an eye on in a group of five. Take one each, or divide them however suits you. They mark where attention goes, so treat the work as shared and rotate the keyboard freely. Each maps to one rubric dimension, so make sure none is left unwatched. An experienced group can ignore these entirely, provided that the building does not crowd out the interrogating.

Printing is optional, so print one set per group if it helps, and cut along the rules. At four people, the Convenor doubles as Steward. At six, add a second Sceptic, since more red-teaming does no harm.

The mapping (canonical):

Angle (handle)	Rubric dimension	In one line
Convenor	Project Definition	Keeps the real problem and the clock in focus
Driver	Technology Stack	Operates the tools and owns the tool choices
Steward	Data Security and Ethics	Guards the red lines
Reporter	Financial and Scalability	Owens the record, the submission and the spoken insight
Sceptic	Human-in-the-Loop Protocol	Tries to break the output and curates the caught errors



Convenor • *Project Definition*

Your angle is that the group stays on a real problem, and on time.

- **In the fifteen minutes.** Settle on your one problem in the first minute, either the seed your group formed around or a real one someone brought, and say how you would know the tool had helped. Share the five angles out, keep half an eye on the clock and make room for everyone to speak.

✂ —————
—————

Reporter · *Financial and Scalability*

Your angle is the record, and the insight the room hears.

- **In the fifteen minutes.** Keep the group's note in the app as you go, or in the rubric template if you are on the paper fallback. Form a view on whether the approach survives the free tier and a larger cohort. Submit the note before the lightning round at 13:33 where you can, and during it at the latest.
- **At the pivot, about 13:26.** Glance at the note and check it is being filled in as you go, instead of leaving it to the last minute.
- **Worth asking.** 'Does this still work next month, for fifty people, for free?'

✂ —————
—————

Sceptic · *Human-in-the-Loop Protocol*

Your angle is testing the output, and deciding where the human has to stay.

- **In the fifteen minutes.** Press on every result, looking for invented detail, bias and confident error, and keep the good catches for the museum. Name the checkpoints the group would want and the things it would never delegate. If there is time, sketch the automation and steering map, and decide whether your oversight ran throughout or at set points.
- **At the pivot, about 13:26.** This is the one moment to speak: bring your best caught error to the group.
- **Worth asking.** 'Where, exactly, does a human have to check before we trust this?'

✂ —————
—————

A personal selection, made in the facilitator's own capacity. It is not the position of the University of Oxford (the facilitator's employer) or of the host, and it is not legal advice.

One-page data decision aid: can I paste this in?

Print one per group, or keep it open in a tab, and use it before anything goes into a free, third-party AI tool. It is a thinking aid and not legal advice. Institutional policy, the UK GDPR and research ethics still govern.

The one-line test

If you would not pin it to a public noticeboard in the building, do not paste it into a free consumer AI tool.

Free tiers commonly reuse your input to train models, and settings change without notice. Treat anything you type as potentially non-private and non-retractable.

Walk the four questions

1. Is it personal data? Does it relate to an identifiable living person (names, emails, student or participant records, anything that could single someone out, even indirectly)? If yes, do not paste it as it is. Anonymise or synthesise it, and it is then no longer personal data. If it has to stay identifiable, a lawful basis is necessary but nowhere near sufficient: you also need a tool your institution has cleared, with a processor contract behind it, which a free consumer tool is not.
2. Is it special category data under Article 9 of the UK GDPR? That means racial or ethnic origin, political opinions, religious or philosophical beliefs, trade union membership, genetic data, biometric data used to identify someone, health, sex life or sexual orientation. Criminal offence data, under Article 10, carries similar weight. If yes, stop there, because a free consumer tool is the wrong place, and your options are to anonymise beyond recognition, to synthesise or to abstain. Children's data and commercially sensitive material are not special category, but both carry extra duties, so treat them the same way here.

3. Is it confidential, unpublished or restricted? This covers unpublished data or manuscripts, grant drafts, peer-review files, material covered by a non-disclosure agreement, anything under embargo and anything shared without co-author or participant consent. If yes, stop there, because confidentiality and embargoes are not yours alone to waive.
4. Do you hold the rights? Is the text or image yours to share, or is it someone else's intellectual property? If unsure, treat it as restricted.

If you cleared all four, you are probably fine. If any gave you pause, that pause is the signal, so take one of the routes below before you paste anything.

At a glance: fine, careful, never

✓ Generally fine	▲ Pause and treat with care	✗ Do not paste into a free tool
Already-public text (your published abstract)	Lightly disguised real material	Personal data of identifiable people
Fully synthetic or fabricated examples	Aggregated or partially de-identified data	Special category, criminal offence or children's data
Generic, non-identifying questions	Draft prose with no sensitive content	Unpublished data, grants, peer review
Your own writing, where you hold the rights	Anything you would hesitate to email widely	Confidential, embargoed or NDA-covered material

If you cannot paste it: what to do instead

- Anonymise. Remove names, identifiers and anything that singles someone out, and check that it cannot be re-identified by combination.
- Synthesise. Recreate the structure of the problem with invented content. A critique of a method works just as well on a paraphrased method.
- Abstract up a level. Discuss the shape of the task, and leave out the sensitive specifics.
- Use an appropriately governed tool. An enterprise or paid service under a proper data-processing agreement may be acceptable where a free consumer tool is not. That is a decision for your institution, and it is out of scope today.

- Abstain. Sometimes the right answer is to leave the tool alone. That is a valid finding, and a good one for your museum of caught errors.
-

Disclaimer: these materials and the tools they reference are the facilitator's own choices, made in a personal capacity, and do not represent the views or official position of the University of Oxford, the facilitator's employer, or of the host, the University of Westminster. The University of Oxford accepts no liability for the selection, use or outcomes of any third-party tool. Participants remain solely responsible for compliance with their own institutional policy, the UK GDPR and research ethics. This wording is a template and is not legal advice.

Evaluation rubric and note template

This is where the group records its work through the day. By default, you capture your work in the workshop app (genai-rt.web.app). Its short form covers the essentials: the problem, the artefact, the errors you caught, the automation–steering map, the oversight model, your insight and a field reflection. Nothing needs setting up beyond the passcode. The fuller template below adds the five rubric dimensions, each scored from 1 to 5 (1 Nascent, 3 Developing, 5 Robust), a deliverable link, the museum table and the societal reflection. It is the note for the HackMD or paper fallback, and for a group that wants to record its scores and reflection after submitting in the app.

The rubric is a thinking tool and a shared record, and once it is archived, it forms part of the open account of what this cohort learned.

The format is deliberately short, because the build window in Part 2 is about fifteen minutes. Get the core three first: the artefact (your tool run on your problem), one caught error and one insight for the lightning round. If you have about five minutes left, add the automation–steering map and say whether your oversight was interwoven or staged. The rubric scores and the societal reflection come later, when you tidy the note. A thoughtful partial note is worth more than a rushed full one.

A. The HackMD fallback (about five minutes, led by the Reporter, and only if you cannot use the app)

1. Go to hackmd.io and sign in. It is free, and you can use a Google, GitHub or email login. There is nothing to install.
2. Click '+ New note'. The left pane is Markdown and the right pane shows a live preview. You can ignore the formatting and type straight into the headings.
3. Copy the whole template in section B below (everything inside the box) and paste it into the note.
4. Fill it in as you work, and avoid leaving it all to the end. Rough notes are fine.
5. Make it link-shareable: open the note's sharing or permissions control and set it so that anyone with the link can read (you do not need to grant editing).

6. Copy the note's link.
7. Share that link with the facilitator before the lightning round at 13:33 where you can, and during it at the latest. Show the note on screen with your table number and track. You do not need a GitHub account, because the facilitator gathers every group's link.

If HackMD will not let you sign up (some logins are restricted), there is no need to lose time on troubleshooting, so switch to paper straight away. Write the same template in any shared document or on paper. It is plain Markdown, and the facilitator can transcribe one note into the archive afterwards.

After the session, the facilitator exports every note whose group opted in to sharing, removes anything that should not be public and commits it under `submissions/`.

B. The note template (paste this box into HackMD or copy it on paper; the app has its own shorter form)

```
# [Group NN] · Track [A / B / C / D] · 9 September 2026

Roles taken (no names needed): Convenor · Reporter · Driver · Sceptic
· Steward
The real problem we brought: ...
Share publicly? yes / no (yes means the group is content for this note,
which should name no one, to enter the public archive)

## 1. Project Definition
- The real problem, in one sentence a colleague would recognise: ...
- Success criterion (how we would know the tool actually helped): ...

## 2. Technology Stack
- Tool(s) used, and where they sit on the spectrum (off-the-shelf / no-
code / IDE-API): ...
- Why this level and not one up or down (control, data exposure,
effort): ...

## 3. Data Security and Ethics
- What we put into the tool, and how we de-identified it: ...
- Red lines we held; UK GDPR, ethics, disclosure of AI use, IP, consent
```

and fairness: ...

4. Financial and Scalability Constraints

- Free-tier limits we hit or foresee: ...
- What breaks first at scale (cost, usage limits or trust): ...

5. Human-in-the-Loop Protocol

- Automation-steering map. Break the project into its phases; for each, note what is automated (the tool does it) and what is human-steered (we decide, verify or override):

Phase	Automated (AI does)	Human steering (we decide / verify / override)
---	---	---
...	...	...

- Interwoven or staged? Is our oversight interwoven (continuous, with a human in the loop at every step) or staged (concentrated at distinct checkpoints between phases)? Which did we choose, and why?
- What our choice costs, and how we offset it. Interwoven oversight is thorough but heavy, and checking every step can dull attention until things are waved through out of habit. Staged oversight is lighter, but errors can accumulate unnoticed between checkpoints: ...
- Checkpoints we set in advance: ...
- What we will never delegate to the tool: ...
- Who is accountable for the final judgement: ...

Deliverable

- Link or screenshot of the artefact we made: ...

Museum of caught errors

What happened	How we caught it	What it signals about where humans must stay
---	---	---
...	...	...

Societal reflection (complete when you tidy and share the note; no score, just think)

- Thinking versus writing: if the tool did the writing, what thinking did we still have to do, and what might we lose by offloading it? ...
- Fairness: could this widen or narrow disparities (first versus second language, more versus less support, career stage, socioeconomic background)? Who gains, and who is left behind? ...
- Disclosure: would we disclose using AI here? Would we, or should we, disclose the conditions it offset (for example, writing in a second language or limited support)? Why or why not? ...

- Disciplinary norms: is our field generally under- or over-using generative AI for this kind of task? Why, with what consequences and how would we redress it over the next two years? ...

Lightning-round insight (up to 45 seconds, no slides)

> Pick the strongest of three: the single most significant limitation we found; the most

> important human-in-the-loop safeguard we built in; or one honest thing this exposed

> about how our field is over- or under-using AI for this task:

>

> ...

Self-assessment (complete when you tidy and share the note; 1 Nascent · 3 Developing · 5 Robust)

Dimension	Score (1–5)	One-line justification
-----------	-------------	------------------------

Project Definition		
--------------------	--	--

Technology Stack		
------------------	--	--

Data Security and Ethics		
--------------------------	--	--

Financial and Scalability		
---------------------------	--	--

Human-in-the-Loop		
-------------------	--	--

C. The rubric anchors

Use these to place yourselves honestly. Most groups will sit around 3 on most dimensions in a short applied session, and that is as it should be. A thoughtful 3 with a sharp caught error is worth more than an unexamined 5. Use 2 and 4 for work that sits between two anchors: 2 is nearer Nascent than Developing, and 4 is nearer Robust than Developing.

1 · Project Definition

Is the real problem, and its success criterion, clear?

- **1, Nascent:** the task is vague or a toy ('use AI on something'), with no success criterion.
- **3, Developing:** a real problem is named and there is a rough sense of what 'better' would look like, although the scope may be broad.
- **5, Robust:** a specific, bounded, genuinely held problem a colleague would recognise, with an explicit success criterion you could actually test.

2 · Technology Stack

Is it the right tool on the spectrum, and why?

- **1, Nascent:** one tool used by default, with no awareness of alternatives or of the spectrum.
- **3, Developing:** the choice is explained, with some sense of where it sits and what moving a level up or down would change.
- **5, Robust:** deliberate placement on the spectrum, with a reasoned trade-off of control, data exposure and effort, and alternatives weighed.

3 · Data Security and Ethics

Were the red lines held, and were the UK GDPR, ethics and fairness addressed?

- **1, Nascent:** personal, special-category or confidential data used without thought, and no consideration of disclosure, consent or fairness.
- **3, Developing:** red lines mostly respected and data de-identified, with some ethical reflection on disclosure, IP, bias or fairness.
- **5, Robust:** red lines explicitly held; anonymisation or lawful basis reasoned; disclosure, IP, consent, bias and fairness addressed; and a note that would withstand being public.

4 · Financial and Scalability Constraints

Does the free tier hold, and what breaks at scale?

- **1, Nascent:** no thought to cost or scale, and an assumption that it stays free forever.
- **3, Developing:** free-tier limits noted, with a rough sense of what scaling would cost or require.
- **5, Robust:** clear about free-tier viability, lock-in and the precise point at which data, cost or trust breaks the approach.

5 · Human-in-the-Loop Protocol

Are there explicit checkpoints, verification and accountability, and a deliberate map of automation against human steering?

- **1, Nascent:** output trusted as it is, with no checkpoint, no named accountability and no sense of where automation ends and human steering begins.

- **3, Developing:** some verification done, one or two checkpoints identified and an implicit split between the automated and human-steered steps.
 - **5, Robust:** a deliberate automation–steering map; oversight consciously chosen as interwoven or staged, with the weakness of that choice offset; an explicit list of what is never delegated; and a named person who owns the final judgement on correctness, ethics and attribution.
-

Disclaimer

These materials and the tools they reference are the facilitator's own choices, made in a personal capacity. They do not represent the views or official position of the University of Oxford, the facilitator's employer, or of the host, the University of Westminster. The University of Oxford accepts no liability for the selection, use or outcomes of any third-party tool. Participants remain solely responsible for compliance with their own institutional policy, the UK GDPR and research ethics.

This wording is a template and is not legal advice.

Worked examples: one per track

Use your group's seed, or a real, non-confidential problem a member brings, and turn to the example below only if neither is to hand. It is also what you use when the Wi-Fi fails and you switch to critiquing on paper.

The material below is fully synthetic and safe to paste, with no real people, data or unpublished work, so a group with nothing to hand can start at once. If you are using it, still spend a minute on the data decision aid: say what you would have had to strip out had this been your own material, and record that under Data Security and Ethics. Each example is sized for the short build window and leads naturally to a caught error.

Track A · Methodological Blind-Spot Detector

Scenario

You are about to submit the study below. Use an AI tool as a critical reviewer, then verify which of its objections are real, which are generic and which are wrong.

Synthetic methods snippet (safe to paste): 'We ran a cross-sectional online survey of 200 marketing managers recruited through a LinkedIn post. Each manager completed a 12-item self-report scale of "authentic leadership" and, in the same questionnaire, rated their own team's "innovativeness" on a 5-item scale. Authentic leadership correlated positively with team innovativeness ($r = .46$, $p < .001$). We conclude that authentic leadership causes higher team innovation and recommend leadership training to raise innovation.'

First move

Ask it openly what is wrong with the design, without naming the lenses, because a prompt that lists the failure modes also hides them. Classify each point it raises as real, generic or wrong. Only then run the second pass across explicit lenses (sampling and recruitment, common-method bias, construct validity, causal inference and generalisability), and note what the guardrails added.

What to watch for

A strong tool should flag the causal claim from cross-sectional data, common-method or single-source bias (the same person rates both variables), self-selection via LinkedIn and self-report of one's own team. Note whether it also pads the list with vague, always-true objections, and whether it ever invents a statistic or a citation. Those are your museum pieces.

Track B · Accessible Executive-Function Layer

Scenario

Turn a messy set of meeting notes into a clean, owned action list, then check that the tool did not quietly invent or misattribute anything.

Synthetic meeting notes (safe to paste; roles only, no real people): 'Centre meeting. Talked about the seminar series. Someone should sort speakers for the autumn, maybe the deputy director? Budget underspend, need to use it before July or we lose it. The new PhD reps raised workload worries. Website is out of date (old staff list). Agreed to revisit the mentoring scheme but no one said who. Ethics turnaround is too slow. Chase the committee. Next meeting in three weeks, same room hopefully.'

First move

Ask for a structured action list with owner, deadline and priority, plus a separate 'unclear or needs an owner' list. Build it as a reusable template your group could run again on the next set of notes.

What to watch for

Does the tool invent owners or dates that the notes never stated ('Deputy Director, by 30 June')? Does it silently drop the items with no owner instead of flagging them? Does it confidently over-formalise a vague discussion? The safeguard you design ('every action must trace to a line in the notes; no invented owners') is the deliverable.

Track C · Rapid Prototyping for Knowledge Translation (technical stretch)

Scenario

Build a tiny interactive that communicates a finding, and verify that it is actually correct before you would ever show it to anyone.

Synthetic finding and data (safe to paste): 'In a (made-up) study, average reading time for a one-page brief fell as font size rose, then rose again when the font got too large. Synthetic data, mean seconds: 10 pt → 95 s; 12 pt → 78 s; 14 pt → 70 s; 16 pt → 72 s; 18 pt → 81 s.'

First move

Ask a tool with a preview (for example, Claude artifacts or Gemini's Canvas) to build a small, self-contained HTML page with a labelled bar or line chart of these values, plus a one-sentence plain-language takeaway. Keep a verification log as you go.

What to watch for

Do the plotted numbers match the data exactly? Is the takeaway honest about the U-shape (fastest reading around 14 pt), or does it over-simplify to 'bigger is better'? Check colour contrast and that the axes are labelled. Then add a sixth point of your own, 20 pt, and see whether the axis, the scale and the takeaway update correctly or whether the tool hard-coded them. The fact that it rendered does not mean that it is correct, so log every gap.

Track D · Public Engagement Translator

Scenario

Translate a hedged, cautious finding for three audiences, then audit each version for accuracy drift.

Synthetic abstract (safe to paste): 'In a small, preliminary, correlational study of 60 undergraduates, students who reported using a structured note-taking app tended to report slightly higher exam confidence (not exam marks). The effect was modest, the sample was not representative, and no causal claim can be made. Replication is needed.'

First move

Produce three versions (an interested public, a policy reader and secondary-school pupils), then line each up against the source and mark every over-claim, dropped caveat or false certainty.

What to watch for

Watch the tool quietly upgrade 'tended to report slightly higher confidence' into 'boosts exam results', drop 'small, preliminary, correlational' and add a confident headline the evidence cannot bear. Your fidelity checklist ('every claim traceable to the source; all caveats survive; no causal language') is the artefact.

Disclaimer: these materials and the tools they reference are the facilitator's own choices, made in a personal capacity, and do not represent the views or official position of the University of Oxford, the facilitator's employer, or of the host, the University of Westminster. The University of Oxford accepts no liability for the selection, use or outcomes of any third-party tool. Participants remain solely responsible for compliance with their own institutional policy, the UK GDPR and research ethics. This wording is a template and is not legal advice.

Starter-prompt library

Each prompt below is a starting point. Adapt freely, replace the bracketed parts and remember the data decision aid before you paste anything.

Each track has two prompts, and the order matters. Run the first pass as it stands, without the guardrails, because a prompt that forbids a failure mode also hides it, and the failure is what you are here to catch. Then design the second prompt as the safeguard your group would actually adopt, and note the difference between them. That difference is your finding.

Prompting for friction: five habits

1. Make it disagree with itself. Ask for the strongest case against its own answer.
 2. Ask it to sort its own claims. 'Mark each claim as confident, uncertain or a guess, and say what would change your mind.' These labels are the model's guess about itself, with no measurement behind them, and a confidently wrong claim will be labelled confident. Use them only to decide what to verify first, then verify it elsewhere.
 3. Demand checkable sources. 'Cite sources I can verify. If you are not sure a source exists, say so instead of inventing one.'
 4. Ask what it is missing. 'What would an expert say you have overlooked?'
 5. Keep the human in the loop. 'List what a human must check before this is used.'
-

Track A · Methodological Blind-Spot Detector

First pass, unguarded

Naming the flaws for it would tell you only that it can follow a list, so leave them unnamed.

'Act as a sceptical peer reviewer. Here is a study design: [PASTE ANONYMISED DESIGN]. What is wrong with it? For each issue, say how serious it is and why.'

Then classify every objection yourselves as real, generic or wrong, and check any factual claim outside the tool. Only now add the guardrails, as the protocol you would give a colleague.

‘Review it again across these lenses: sampling and recruitment, common-method bias, construct validity, causal inference, generalisability and ethics. Separate study-specific weaknesses from objections that would apply to almost any study, and do not invent citations.’

‘Now argue the opposite: what is methodologically strong here, and which of your own criticisms is weakest?’

Track B · Accessible Executive-Function Layer

First pass, unguarded

Watch what it does with the items that have no owner and no date.

‘Turn these meeting notes into a structured action list with columns for action, owner, deadline and priority: [PASTE REDACTED NOTES].’

Check every owner and date against the notes. Then write the safeguard, which is the deliverable for this track.

‘Do it again. Only include owners and dates that are explicitly stated, put anything vague in a separate list headed “needs an owner or unclear”, and do not invent details.’

‘Review your own output: did you add any action, owner or date that was not in my notes? List exactly what you inferred and what I gave you.’

Track C · Rapid Prototyping for Knowledge Translation

‘Build a single self-contained HTML page that shows this data as a labelled chart, with a one-sentence plain-language takeaway: [PASTE SYNTHETIC DATA]. Use accessible colours and label the axes. Add a short comment noting any assumption you made.’

‘Now act as a reviewer of your own code: where could the numbers be wrong, what breaks with unexpected input, and what must a human verify before this is published?’

Track D · Public Engagement Translator

First pass, unguarded

Telling it to keep the caveats would hide the drift you are looking for.

‘Rewrite this abstract for three audiences, an interested general reader, a policy reader and secondary-school pupils: [PASTE PUBLIC ABSTRACT].’

Line each version up against the source yourselves and mark every over-claim, dropped caveat and added certainty. Then write the checklist as a prompt.

‘Do it again, preserving every caveat and uncertainty, and adding no causal claim the abstract does not make.’

‘Audit your three versions against the original. List every place you strengthened a claim, dropped a caveat or added certainty the source did not have.’

General critical-use prompts (any track)

'Before answering, tell me what you would need to know to answer well, and what assumptions you are making.'

'Give your answer, then a section headed "Where I might be wrong" and a section headed "What a human should check".'

'I think your answer is wrong because [REASON]. Defend it or revise it, and tell me which.'

'Map the task: list its phases, mark what you would automate and where a human must steer or verify, then say whether oversight should be continuous (interwoven) or at checkpoints (staged), and why.'

Disclaimer: these materials and the tools they reference are the facilitator's own choices, made in a personal capacity, and do not represent the views or official position of the University of Oxford, the facilitator's employer, or of the host, the University of Westminster. The University of Oxford accepts no liability for the selection, use or outcomes of any third-party tool. Participants remain solely responsible for compliance with their own institutional policy, the UK GDPR and research ethics. This wording is a template and is not legal advice.

Group quick-start one-pager

Copies	20 copies
Sides	Single-sided
Paper	A4 portrait
Finishing	Loose, two laid on each table
Length	1 page

Two per table, on top of the booklet.

Your group's quick start

One or two copies per table. Everything here works on a phone, though a laptop helps if you pick Track C and want to build or visualise something.

1. Open the app

genai-rt.web.app

Type that into any browser, or scan the QR code on screen. No account and nothing to install.

2. Start your group

One person sets it up. The same form starts a new group or joins an existing one.

- To start a new group, leave the join code blank, choose a group name that does not identify anyone (such as 'Otters' or 'Team Kelp'), type the passcode the facilitator reads out and press Continue.
- The app then shows a six-character join code. Read out the group name and the code, since anyone joining needs both, and they leave the passcode blank. One shared device is fine too.
- To join a group that already exists, open the same address, type the same group name and enter the code. No passcode is needed to join.

If the passcode is refused, check it with the facilitator. If the name is taken, pick another.

3. Choose a scenario

Pick your track from the dropdown.

- A. Methodological Blind-Spot Detector
- B. Executive-Function Layer
- C. Rapid Prototyping
- D. Public Engagement
- Own problem. A real one a member brings. Anonymise it or synthesise it, and abstain if you are in any doubt. Never use personal, student, unpublished or confidential data.

4. The core three

This is the heart of the fifteen minutes. If this is all you manage, you have done the task.

1. One artefact. Run your tool on your problem and produce something real, such as a critique, a template or a translation.
2. One or two caught errors. Note the moments where it was fluent but wrong, and keep them. This is your museum of caught errors.

3. One insight. A limitation, a safeguard or the honest thing this exposed. This is your line for the lightning round, which allows up to 45 seconds per group.

With about five minutes to spare, go further.

- Automation—steering map. Which steps the tool ran, and where you steered by deciding, verifying or overriding.
- Oversight model. Interwoven means checking throughout, and staged means checking at set points. Say which you chose and what it cost.
- Field reflection. Is your field over- or under-using generative AI here, with what consequences and how would you redress it?

5. Share and submit

- Sharing is optional. Tick the consent box if your group is happy for its non-identifying note to enter the public archive on GitHub. Unticked, it stays out of the archive, although approved work still shows on the session's passcode-gated dashboard.
- Submitting. Choose a scenario first, then submit when you are done, before the lightning round at 13:33 where you can, and during it at the latest.
- After submitting. The form becomes read-only while the facilitator reviews it. If they reopen it with a note, you can edit and resubmit.

Good to know

- Your work is safe. The app saves as you go. If you reload the page or lose Wi-Fi, your note stays on the device and syncs when you reconnect.
- Data safety. No personal or confidential data goes into a free AI tool. When in doubt, anonymise, synthesise or abstain. The only identifier the app holds is your group name.
- The timer. A chip in the corner shows the time left. It is advisory and never locks you out.
- Roles are optional. The Convenor keeps the problem and the clock, the Driver runs the keyboard, the Steward guards the red lines, the Reporter owns the note and the spoken insight and the Sceptic tries to break the output.
- If the app will not load, go to hackmd.io and use section A of the rubric in your pack, or the paper copy the facilitator gives you. The headings are the same.

A personal selection, made in the facilitator's own capacity. It is not the position of the University of

Oxford (the facilitator's employer) or of the host, and it is not legal advice.

Run sheet and morning checklist

Copies	1 copy
Sides	Double-sided
Paper	A4 portrait
Finishing	Stapled
Length	6 pages

Kept with the cue cards.

Facilitator run sheet

The single-glance timeline for the day. It is operational only: the reasoning is in the [facilitator guide](#), and the holdable prompts are on the [cue cards](#). Slide numbers refer to [slides.md](#).

A featured workshop in two 30-minute parts, split by a networking lunch.

Before Part 1 (see the [morning checklist](#))

When	Action
11:45	Projector on, deck open in presenter mode, Wi-Fi and power confirmed.
11:50	Create and pin the fallback-links issue, so that it is there if a group fails over to HackMD later (participants do not post to it).
11:52	Pre-flight the app (genai-rt.web.app): sign in to the private dashboard and confirm that it loads. The full reset, meaning today's passcode, the timer and the clearing of rehearsal groups, happens in the lunch break (see the day-of reset). No pre-named notes are needed.
11:55	Lay out one group pack booklet and one or two group one-pagers on each table of five, with spare pens and paper. Post the ten seed ideas around the room, each with its track letter, ready for the 12:26 formation.
11:58	Visible timer ready, the lightning-round tally sheet and a pen to hand, title slide up (slide 1). Start at 12:00.

Part 1 · Conceptualising the use of AI in research · 12:00–12:30

Clock	Segment	Slides	Facilitator does
12:00	Welcome and stance	1–5	Start on time. Disclaimer aloud. Discernment, not enthusiasm; the aim is insight, and all are welcome.
12:03	Saved or burned	6	A sentence each with a neighbour, two minutes, then move on.
12:05	A long line of tools	7	AI in the history of tools and the productivity economy: the question is how, with discernment.
12:08	Getting started	8	What these tools are, pick one free tool, how to ask. For all levels.
12:10	Friction is a signal	9	The central idea. Lee et al. and Drosos et al. in a sentence each.
12:14	Spectrum, red lines, ethics	10–12	One breath each. Name the teaching angle (students' data, assessment, disclosure).
12:21	The bigger questions	13	Paradigm shift, fairness, disclosure, posed as live questions for Part 2.
12:23	Human in the loop and the map	14	Accountability stays human; interwoven or staged.
12:26	What Part 2 asks, then form groups	15–16	Two minutes on the rubric, two on formation. Groups of five form at a seed, with the seed letters posted around the room beforehand so that people walk to a sign. Each group's track follows from its seed, so point it to the matching worked example and starter prompts in the pack.

Networking lunch · 12:30–13:15

When	Action
12:30	Not your session, but a networking break before Part 2. Leave slide 17 up throughout: it is the only thing telling people to come back to the same table.
12:45–13:10	Reset the app: today's passcode, the timer length and any rehearsal groups to clear. See the day-of reset .
13:12	Slides 18 and 19 (the Part 2 divider, then the QR code and the address) on screen before anyone sits down, and today's passcode written up where the room can read it. Groups can then start joining as they sit, instead of waiting for you to read it out.

Part 2 · Practical AI for research and teaching · 13:15–13:45

Clock	Segment	Slides	Facilitator does
13:15	Settle in and orient	19–24	Re-find groups; agree the one problem (60 seconds); app open (genai-rt.web.app); read out the passcode; red lines on. Slides 21–24 go past at pace: point at them, do not present them, since every word is in the pack.
13:18	Apply it	25	Press Start countdown as slide 25 goes up. If the settle-in has run late, set Minutes so the countdown still ends at 13:32 rather than fifteen from now. Circulate. Stop sensitive data, unstick stalls, one provocation each. Where a group has time, prompt the map (interwoven or staged).
about 13:26	Pivot	26	Thirty seconds, keyboards down. Stop building, start interrogating.
13:33	Lightning round	27	Visible timer. 45 seconds each with eight groups; announce 35 seconds if there are nine or ten, since handovers cost five to ten seconds apiece. Time the round and jot one line per group on the tally sheet. Leave every approval until after the session, since reading a note properly takes longer than the close allows (collect a HackMD link only from a fallback group).
13:40	Close	28–34	Synthesise two or three threads. Take-home actions. Reading. Follow-up. Close at 13:45.

After

When	Action
Same day	Check that every app group is approved on the dashboard, and that any HackMD-fallback group's link is in the issue. Nudge any that are missing.
+1–2 days	Export the approved, consented work from the dashboard (Export approved (Markdown)), export any HackMD-fallback note too, and archive under <code>submissions/</code> (see submissions/README.md).
+3–5 days	Send the follow-up email with the responsible-use commitment and the archive link.

A personal selection, made in the facilitator's own capacity. It is not the position of the University of Oxford (the facilitator's employer) or of the host, and it is not legal advice.

Morning checklist

Tick these before the room fills. The session is low-tech by design, so most of the risk sits in the first ten minutes.

Technology

- ☐ Projector or screen working, with the laptop output tested at the room's resolution.
- ☐ Deck open: the committed `handouts/slides.pdf`, or `dist/slides.html` in a browser for presenter view, which `npm run build` generates. Keep a second copy on a phone or a memory stick, since a dead laptop otherwise takes the deck with it.
- ☐ Wi-Fi confirmed on a guest device, on the network participants will actually use.
- ☐ Power: enough sockets and extension leads for tables of five.
- ☐ `genai-rt.web.app` (the workshop app) opens on a guest device and shows the login form, which is the proof that its code loaded rather than merely the page. The facilitator's Google sign-in to the dashboard works.
- ☐ A chat tool (ChatGPT, Claude, Gemini or Copilot) reachable on the guest network.
- ☐ A visible timer (phone, slide timer or kitchen timer) for the segment timings and the lightning round.

App pre-flight (facilitator)

- ☐ Dashboard reachable, and Google sign-in completes with pop-ups allowed.
- ☐ Dashboard clear of any previous session's submissions. Clearing is console-only, so if anything is there, note it and clear it during the lunch reset.
- ☐ Today's passcode and the timer are set during the lunch break, so leave them until then. See the [day-of reset](#) for the full list.

Printed materials (per table of five)

- ☐ Group pack booklet, one per table: cover, role cards, data decision aid, rubric, worked examples and starter prompts.
- ☐ Group one-pager, one or two per table.
- ☐ Seed-idea signs, ten of them, one per page in A4 landscape, printed and posted around the room before 12:00. Each carries its track letter, and they are generated from the quick group seeds in `project_tracks.md`.
- ☐ Table numbers, one per table, large enough to read across the room. Groups form by moving to a seed idea at the end of Part 1 and then have to find the same table after a 45-minute lunch, and the table number is what they write down.
- ☐ Spare pens and paper. Paper is the first fallback for a group with no device, and HackMD the second for a group that has one but cannot reach the app.

Facilitator kit

- ☐ Cue cards printed and in order.
- ☐ Run sheet ready.
- ☐ Lightning-round tally printed, with a pen, for the one line per group that the close draws its threads from.
- ☐ Water, a clock you can see and a way to get the room's attention.

Details to confirm

- ☐ The name, date, venue, times and affiliation in the deck and the emails are yours (they are set to the facilitator's own: Department of Education, University of Oxford, and the University of Westminster on 9 September 2026).
 - ☐ The placeholders that remain were filled before sending: `[VENUE]` and `[POLL URL]` in the pre-workshop email, and `[FEEDBACK URL]` and `[SURVEY SUMMARY]` in the follow-up email.
 - ☐ The University of Oxford crest is embedded in `themes/workshop.css` (its source is `assets/logo.svg`), and the deck was rebuilt after any change.
-

A personal selection, made in the facilitator's own capacity. It is not the position of the University of Oxford (the facilitator's employer) or of the host, and it is not legal advice.

Lightning-round tally

Copies	1 copy
Sides	Single-sided
Paper	A4 portrait
Finishing	Loose, on a clipboard or the lectern
Length	1 page

Somewhere to write one line per group during the round at 13:33, which the close at 13:40 then draws its threads from.

Lightning round: one line per group

Somewhere to write while the room talks. The close at 13:40 asks for two or three threads drawn from the round, and ten contributions of 45 seconds each are more than anyone recalls unaided. Fill in the middle columns as each group speaks, and put a word in the last one for the thread it belongs under.

#	Group	Track	Their line	Thread
1				
2				
3				
4				
5				
6				
7				
8				
9				
10				
11				

The threads to name at 13:40

Write them here as they emerge, then read them back to the room before connecting them to the day: friction as a signal, the spectrum of tools, the red lines and the human in the loop.

1	
2	
3	

A personal selection, made in the facilitator's own capacity. It is not the position of the University of Oxford (the facilitator's employer) or of the host, and it is not legal advice.

Day-of app reset

Copies	1 copy
Sides	Single-sided
Paper	A4 portrait
Finishing	Loose
Length	2 pages

| The passcode, timer and export steps for the lunch break.

App reset checklist for the day

This checklist is operational only. The [facilitator guide](#) holds the rationale, the [run sheet](#) the timeline and the [app README](#) the full deployment steps.

A short reset, in the minutes before Part 2, that gets the live app into a clean, ready state for Westminster on 9 September 2026. The app ([genai-rt.web.app](#)) is where groups capture their work by default, and HackMD or paper is the fallback.

During the lunch break (about 12:45–13:10)

- ☐ Confirm the live app is running the current code. If anything under `firebase-app/` has changed since the last deploy, run `firebase deploy --only firestore:rules,hosting`, which has to send both together: the join rule and the join code the pages write were changed as a pair, so deploying one without the other stops groups joining. Then create a test group on one device and join it from a second before you clear the rehearsal groups.
- ☐ Sign in to the facilitator dashboard at [genai-rt.web.app/facilitator.html](#) with your facilitator Google account, the one named in `firebase-app/firestore.rules` ('Sign in with Google'). Group devices do not sign in here.
- ☐ Confirm that sign-in actually completes and that you are on the dashboard, not on the sign-in card. If it bounces, the project's Google sign-in needs checking in the Firebase console, which is a pre-deployment concern to fix before the room fills.
- ☐ Clear the rehearsal and test groups. Deletion is console-only, with no button on the dashboard. In the Firestore console, delete stray test documents from the `groups` collection and the matching `groupNames` entry, or the name stays reserved.
- ☐ Confirm that the facilitator group list is clear of rehearsal entries (that is where draft and submitted test groups show) and that the public board at [genai-rt.web.app/dashboard.html](#) is empty.
- ☐ Set today's passcode in the Session passcode panel and write it where the room can see it (it overwrites any rehearsal value). The status should then read 'Passcode set. Read it out to the room.' Confirm that the public dashboard opens with this passcode too. Until it is set, no group can start.

- ☐ In the Session timer panel, set Minutes (the default is 15, the build window) but do not start it yet. Set it to 14 rather than 15, so the countdown ends at 13:32 and leaves a minute to submit inside the build window. Pressing Start again with a new number simply restarts the countdown, so if the settle-in runs late, set Minutes to the time remaining until 13:32.
- ☐ Have the 'Open the workshop app' slide ready, with the large URL [genai-rt.web.app](#) and the scannable QR code, and put it up at about 13:12 so that groups can start joining as they sit.

At the start of Part 2

- ☐ Read out the passcode and the URL ([genai-rt.web.app](#)), and show the QR slide. 'One device per group creates, then reads out its group name and the six-character code. Everyone else types both and leaves the passcode blank.'
- ☐ Remind groups to choose a group name that does not identify anyone (for example 'Otters' or 'Team Kelp'), which is the one name they type on the login form.
- ☐ Press Start countdown at 13:18, when the 'Fifteen minutes, starting now' slide goes up, and not at the 13:15 settle-in. The corner chip is advisory, a calm countdown that never locks their form. Once it expires, it reads 'Time's up' until you reset it.

During

- ☐ Watch the count line, 'N groups · M awaiting review'.
- ☐ On each submission: Approve, Reopen... (it prompts for a short note) or Rename... as needed.
- ☐ Approve only when the group is done. Approval shows the work on the session dashboard and makes the group's form read-only, because owners can edit only while a note is in draft or reopened. The join code is wiped too, so no new device can join, and Reopen is no longer offered.
- ☐ Rename any identifying name the moment you spot one (Rename..., keeping it short and non-identifying).
- ☐ If a phase overruns, type a new number of Minutes and press Start again, which restarts the countdown at that length.

After Part 2 (once groups have submitted)

- ☐ Export approved (Markdown) from the dashboard is the default path. It exports only approved and consented groups and downloads `YYYY-MM-DD_genai-rt-submissions.md`, dated with the day you export.
- ☐ Check it before committing (names and content), rename it to the workshop date if you exported later, then commit under `submissions/`.
- ☐ Alternatively, from your own machine if pre-configured, run `npm run archive:pr` (it needs `GENAI_RT_PROJECT` and `GENAI_RT_API_KEY` set and `gh` already logged in). To preview, run `npm run archive:pr -- --dry-run`. If either is not ready, use the Markdown export above.

Troubleshooting

- Sign-in fails: allow pop-ups for the site and try again, since the pop-up path is the reliable one. Then confirm that you are using the pre-registered Google account, and check that Google sign-in is enabled for the project in the Firebase console.
- A group cannot start: the passcode is not set, or they have a typo. Re-read it, since the create path needs the session passcode.
- Reopen... or Rename... does nothing: the browser has suppressed this page's dialogs, which it offers to do after several in a row. Reload the dashboard page.
- Timer not showing: glance at your dashboard's Session timer status line first. It reads 'Running: MM:SS left ...' when live and 'Timer is off.' when stopped. The group chip appears while the countdown is running and reads 'Time's up' after it expires.
- A group's page sits on 'Connecting...': after about nine seconds it says so itself and names the fallback, and it says so at once if the code could not be fetched at all. The usual cause is a Wi-Fi sign-in page that has not been accepted on that device, so have them open any other site first, then reload. Failing that, HackMD or paper.
- App wobbles: tell the room to switch to HackMD (hackmd.io), with the same headings. In-progress edits are queued on the device and sync on reconnect, so nothing already typed is lost.

A personal selection, made in the facilitator's own capacity. It is not the position of the University of Oxford

(the facilitator's employer) or of the host, and it is not legal advice.

Part 4 — put up around the room

Seed-idea signs

Copies	1 copy
Sides	SINGLE-SIDED
Paper	A4 LANDSCAPE
Finishing	Heaviest paper available
Length	10 pages

Posted around the room before 12:00, spread out so ten groups can gather without crowding. Do not let the printer rotate or shrink these to fit portrait.

TRACK

A

Stress-test a study design

TRACK

A

**Critique an assessment or
marking rubric**

TRACK

B

**Turn messy notes into owned
actions**

TRACK

B

Tame an email or admin backlog

TRACK

B

Plan a module, project or paper

TRACK

C

**A one-page explainer of a
finding**

TRACK

C

**A quick visualisation or
interactive teaching aid**

TRACK

D

A lay summary or press angle

TRACK

D

**Explain a hard concept for
students**

TRACK

D

A social thread from a paper

Table numbers

Copies	1 copy
Sides	SINGLE-SIDED
Paper	A4 LANDSCAPE
Finishing	Heaviest paper available
Length	10 pages

| One on each table, readable across the room.

1

TABLE

2

TABLE

3

TABLE

4

TABLE

5

TABLE

6

TABLE

7

TABLE

8

TABLE

9

TABLE

10

TABLE