

PRINT JOB

Generative AI in research and teaching Workshop print bundle

Item	Copies	Sides	Paper
Part 1 — cut these			
Facilitator cue cards, for cutting (7 pages)	1 copy	SINGLE-SIDED	A4 portrait
Part 2 — collate into group bundles			
Group pack booklet (5 pages)	10 copies, plus one spare	Double-sided	A4 portrait
Group quick-start one-pager (1 page)	20 copies	Single-sided	A4 portrait
Part 3 — facilitator only			
Run sheet and morning checklist (5 pages)	1 copy	Double-sided	A4 portrait
Day-of app reset (1 page)	1 copy	Single-sided	A4 portrait
Part 4 — put up around the room			
Seed-idea signs (10 pages)	1 copy	SINGLE-SIDED	A4 LANDSCAPE
Table numbers (10 pages)	1 copy	SINGLE-SIDED	A4 LANDSCAPE

This cannot be run as one job with one setting. The copy counts differ per item, the two card sections must print single-sided because they are cut, and the room signs are landscape. Each section below opens with its own instruction slip; the slips are not part of the job and need not be printed.

Part 1 — cut these

Facilitator cue cards, for cutting

Copies	1 copy
Sides	SINGLE-SIDED
Paper	A4 portrait
Finishing	Cut once across the middle of every sheet, then keep in clock order
Length	7 pages

14 cards, two to a sheet, running from 12:00 to 13:40. Each card is numbered in its corner, so the order survives the cutting.

12:00 · WELCOME AND STANCE (Part 1)

- **Say:** 'We are here for discernment, not enthusiasm: a deeper insight into the technology, and no training on any one tool. Whatever your confidence, you are welcome.'
- **Do:** read the disclaimer aloud, and make the point that friction is a signal.
- **Watch:** start on time. This is the conceptual part, so keep it brisk.

Cue card 1 of 14

12:03 · SAVED OR BURNED

- **Say:** 'One sentence each with your neighbour: the last time an AI tool saved you or burned you.'
- **Do:** two minutes, in parallel, then pull the room back.

Cue card 2 of 14

12:05 · A LONG LINE OF TOOLS

- **Say:** 'AI is the latest in a long line of productivity tools, from the calculator and the computer to the spreadsheet and now generative AI. Our institutions reward output, so we need to decide when the tool helps and when it does not.'
- **Do:** show the chain on screen, and take no more than a couple of minutes.

Cue card 3 of 14

12:08 · GETTING STARTED (no experience needed)

- **Say:** 'A chat assistant takes a request and writes back. Pick one free tool. Say who you are, what you want and what a useful response would look like, then push back.'
- **Do:** point to the ready examples in the group pack, and reassure the room that mixed groups are a strength and no one is behind.

Cue card 4 of 14

12:10 • FRICTION IS A SIGNAL

- **Say:** 'When the tool resists you, that friction is information. It marks where your judgement does the work.'
- **Do:** Lee et al. and Drosos et al. in a sentence each. Do not lecture.

Cue card 5 of 14

12:14–12:21 • SPECTRUM • RED LINES • ETHICS

- **Do:** one breath each (slides 10–12). Name the teaching angle: students' data, assessment integrity, disclosure.
- **Watch:** these are reference frames for Part 2, and there is no need to master them now.

Cue card 6 of 14

12:21 • THE BIGGER QUESTIONS AND THE HUMAN IN THE LOOP

- **Say:** 'Zoom out: the shift between thinking and writing; fairness, meaning who is helped and who is left out by language, support, career stage or background; and disclosure. There are no answers today, so notice them in Part 2.'
- **Do:** then the human in the loop and the map (slide 14). Accountability stays human.

Cue card 7 of 14

12:26 • WHAT PART 2 ASKS, THEN FORM GROUPS (end of Part 1)

- **Say:** 'Part 2 puts these ideas to work. Now form groups of five: gather at the idea on screen, and the letter is your track.'
- **Do:** show the three things Part 2 asks (slide 15) and the tracks (slide 16), then send them to lunch. End Part 1 at 12:30.

Cue card 8 of 14

12:30–13:15 · NETWORKING LUNCH (not your session)

- **Do:** a networking break. Use part of it for the app reset if it is not already done. Part 2 starts at 13:15.

Cue card 9 of 14

13:15 · SETTLE IN (Part 2)

- **Say:** 'Re-find your group. Quickly agree one problem: the seed or a non-confidential task someone brought. Open the app at genai-rt.web.app. Here is the passcode. Keep the red lines in mind.'
- **Do:** take about three minutes: regroup, agree the problem in roughly a minute, then open the app and read out the passcode. Protect the fifteen-minute build window.

Cue card 10 of 14

13:18 · APPLY IT (15 min)

- **Let them work.** Stop sensitive data, unstick stalls and offer one provocation each.
- **Push:** a caught error first, then the one line for the room. The optional boxes in the app are for groups with time left.

Cue card 11 of 14

About 13:26 · PIVOT 🕒

- **Say:** 'Thirty seconds, keyboards down. What did it make easier, and at what cost? Where did it resist you? Have you held the red lines? Now switch from building to interrogating.'

Cue card 12 of 14

13:33 · LIGHTNING ROUND (7 min)

- **Say:** 'One insight each and no slides. I will call the time. Give us the biggest limitation, the best safeguard or how your field over- or under-uses AI here.'
- **Say, before they submit:** 'Tick the sharing box at the foot of the form only if you are happy for your note to join the public archive. It is optional.'
- **Do:** visible timer. 45 seconds each with eight groups, or announce 35 seconds if there are nine or ten, since each handover costs five to ten seconds. Still behind after six? Take one line from the rest. Jot each group's thread, and leave every approval until after the session.

Cue card 13 of 14

13:40 · CLOSE (5 min)

- **Synthesise the round** by naming two or three threads from it.
- **Connect:** friction as signal, the spectrum, the red lines, the human in the loop.
- **Give:** the take-home actions, the further reading and the follow-up. Close at 13:45.

Cue card 14 of 14

Part 2 — collate into group bundles

Group pack booklet

Copies	10 copies, plus one spare
Sides	Double-sided
Paper	A4 portrait
Finishing	Stapled, top left
Length	5 pages

One booklet per table. The cover has fill-in fields, so every table needs its own. The footer numbers every page, so a booklet missing a leaf shows up while collating.

Worked examples: one per track

A worked example is a fallback. Start with the seed your group formed around or a member's non-confidential problem. Use the example for your track only when neither is available. It also gives you something to critique on paper if the Wi-Fi fails.

The material below is synthetic and safe to paste. It contains no identifiable people, data or unpublished work, so a group can begin at once. Spend a minute on the data decision aid anyway. Say what you would remove from your own material, then record it under Data Security & Ethics. Each example is small enough to start in the short build window and should lead to a useful caught error.

Track A · Methodological Blind-Spot Detector

Scenario

You are about to submit the study below. Use an AI tool as a critical reviewer, then verify which of its objections are real, which are generic and which are wrong.

Synthetic methods snippet (safe to paste): 'We ran a cross-sectional online survey of 200 marketing managers recruited through a LinkedIn post. Each manager completed a 12-item self-report scale of "authentic leadership" and, in the same questionnaire, rated their own team's "innovativeness" on a 5-item scale. Authentic leadership correlated positively with team innovativeness ($r = .46$, $p < .001$). We conclude that authentic leadership causes higher team innovation and recommend leadership training to raise innovation.'

First move

Ask it openly what is wrong with the design, without naming the lenses, because a prompt that lists the failure modes also hides them. Classify each point it raises as real, generic or wrong. Only then run the second pass across explicit lenses (sampling and recruitment, common-method bias, construct validity, causal inference and generalisability), and note what the guardrails added.

What to watch for

A useful answer should flag the causal claim from cross-sectional data, common-method or single-source bias (the same person rates both variables), self-selection via LinkedIn and self-report of one's own team. Look for vague, always-true objections and invented statistics or citations. Record them as caught errors.

Track B · Accessible Executive-Function Layer

Scenario

Turn a messy set of meeting notes into a clear action list, then check that the tool did not quietly invent or misattribute anything.

Synthetic meeting notes (safe to paste; roles only, no real people): 'Centre meeting. Talked about the seminar series. Someone should sort speakers for the autumn, maybe the deputy director? Budget underspend, need to use it before July or we lose it. The new PhD reps raised workload worries. Website is out of date (old staff list). Agreed to revisit the mentoring scheme but no one said who. Ethics turnaround is too slow. Chase the committee. Next meeting in three weeks, same room hopefully.'

First move

Ask for a structured action list with owner, deadline and priority. Ask it to place unclear items, or items without an owner, in a separate list. Turn the resulting prompt into a template your group could use again.

What to watch for

Does the tool invent owners or dates that the notes never stated ('Deputy Director, by 30 June')? Does it silently drop items with no owner? Does it over-formalise a vague discussion? Use a safeguard such as 'every action must trace to a line in the notes; no invented owners'.

Track C · Rapid Prototyping for Knowledge Translation (*technical stretch*)

Scenario

Build a tiny interactive that communicates a finding, then verify it before sharing it.

Synthetic finding and data (safe to paste): 'In a (made-up) study, average reading time for a one-page brief fell as font size rose, then rose again when the font got too large. Synthetic data, mean seconds: 10 pt → 95 s; 12 pt → 78 s; 14 pt → 70 s; 16 pt → 72 s; 18 pt → 81 s.'

First move

Ask a tool with a preview (for example, Claude artifacts or Gemini's Canvas) to build a small, self-contained HTML page with a labelled bar or line chart of these values, plus a one-sentence plain-language takeaway. Keep a verification log as you go.

What to watch for

Do the plotted numbers match the data exactly? Is the takeaway honest about the U-shape (fastest reading around 14 pt), or does it over-simplify to 'bigger is better'? Check colour contrast and that the axes are labelled. Then add a sixth point of your own, 20 pt, and see whether the axis, the scale and the takeaway update correctly or whether the tool hard-coded them. The fact that it rendered does not mean that it is correct, so log every gap.

Track D · Public Engagement Translator

Scenario

Translate a hedged, cautious finding for three audiences, then audit each version for accuracy drift.

Synthetic abstract (safe to paste): 'In a small, preliminary, correlational study of 60 undergraduates, students who reported using a structured note-taking app tended to report slightly higher exam confidence (not exam marks). The effect was modest, the sample was not representative, and no causal claim can be made. Replication is needed.'

First move

Produce three versions (an interested public, a policy reader and secondary-school pupils), then line each up against the source and mark every over-claim, dropped caveat or false certainty.

What to watch for

Watch the tool quietly upgrade 'tended to report slightly higher confidence' into 'boosts exam results', drop 'small, preliminary, correlational' or add a headline the evidence cannot bear. Create a fidelity checklist: every claim is traceable to the source, all caveats survive and no causal language appears.

Disclaimer: these materials and the tools they reference are the facilitator's own choices, made in a personal capacity, and do not represent the views or official position of the University of Oxford, the facilitator's employer, or of the host, the University of Westminster. The University of Oxford accepts no liability for the selection, use or outcomes of any third-party tool. Participants remain solely responsible for compliance with their own institutional policy, the UK GDPR and research ethics. This wording is a template and is not legal advice.

Fallback note and rubric

Use this page if the app fails or when your group wants to record a fuller reflection after the session. It does not need attention during the fifteen-minute build.

During the build, use the app to record three things: the result of using a tool on your problem, one or two caught errors and the line you will give the room at 13:33. That line may name a limitation, a checkpoint you would require before use or how your field over- or under-uses AI for this task. The material below supports the fallback and the debrief.

If the app will not work

Open a note at hackmd.io, or take the paper sheet from the facilitator, and copy the headings below into it. Make rough notes as you work. They preserve what happened better than a tidy account written at the end. Give the facilitator the link or sheet before the lightning round at 13:33 where possible, or during it at the latest.

If HackMD will not let you sign up, switch to paper at once. The headings work just as well there, and a sign-up is never worth the minutes that troubleshooting it costs.

Group ____ · Track A / B / C / D / Own problem · 9 September 2026

The problem:

The artefact, and the tool that made it:

Errors caught (what happened, how we caught it, where a human must stay):

Our line for the lightning round:

Optional, if minutes remain

Automation-steering map (the phases, what the tool ran, where we steered):

Oversight interwoven or staged, and why:

Field reflection (over- or under-used here, and what we would do about it):

Archive this note publicly? yes / no

The rubric after the session

For a short applied session, 3 is a reasonable result. A sharp caught error tells you more than an unexamined 5. Use 2 and 4 for work that falls between two anchors.

Dimension	1 Nascent	3 Developing	5 Robust
Project definition. Is the problem, and how you would know it helped, clear?	a vague or toy task, with no success criterion	a problem named, with a rough sense of better	bounded and owned, with a criterion you could test
Technology stack. Is this the right tool on the spectrum, and why?	one tool by default, with alternatives unconsidered	the choice explained, with some sense of the spectrum	placed deliberately, weighing control, data exposure and effort
Data security and ethics. Were the red lines held?	sensitive material pasted without thought	red lines mostly held and the data de-identified	red lines held, with disclosure, IP, consent and fairness reasoned
Financial and scalability. Does the free tier hold?	no thought given to cost or scale	limits noted, with a rough sense of what scaling needs	clear on lock-in and on what breaks first, cost, limits or trust
Human in the loop. Are there checkpoints, verification and accountability?	the output trusted as it stands	some verification, and one or two checkpoints named	oversight chosen and its weakness offset, with one named owner

To record the scores, add a line to your note for each dimension with its number and a sentence saying why.

Worth a few minutes afterwards

Say what thinking you still had to do once the tool had done the writing, and what would go if you handed that over as well. Note who gains from working this way and who is left behind, whether by first or second language, support, career stage or background. Set down whether you would disclose the use of AI here, and whether you would also disclose the conditions it offset. Then judge as honestly as you can whether your field is over- or under-using AI for this kind of task, and what you would want done about it over the next two years.

Disclaimer

These materials and the tools they reference are the facilitator's own choices, made in a personal capacity. They do not represent the views or official position of the University of Oxford, the facilitator's employer, or of the host, the University of Westminster. The University of Oxford accepts no liability for the selection, use or outcomes of any third-party tool. Participants remain solely responsible for compliance with their own institutional policy, the UK GDPR and research ethics.

This wording is a template and is not legal advice.

Group quick-start one-pager

Copies	20 copies
Sides	Single-sided
Paper	A4 portrait
Finishing	Loose, two laid on each table
Length	1 page

Two per table, on top of the booklet.

Your fifteen minutes

Two copies per table, and everything here works on a phone.

From 13:18 to 13:33, work on one problem: the seed your group formed around, a member's non-confidential problem or the ready example for your track. You do not need to finish. When the tool resists you, make a note of it. That is often where human judgement matters most.

What the fifteen minutes ask

1. Use a tool on your problem and keep a tangible result: a critique, a template, a chart or a translation. Rough is fine.
2. Note one or two places where the tool was fluent but wrong. Keep the output as it appeared. Those errors show what needs checking.
3. Prepare one line for the lightning round. You have up to 45 seconds and no slides. Name the limitation you found, the checkpoint you would require before use or how your field over- or under-uses AI for this task.

If time is short, write the line for the room before anything else. A partial note with a clear observation is enough. If your group finishes early, use the app to map what the tool automated and where you steered, choose between interwoven and staged oversight, and reflect on your field.

Choose three jobs in the first minute. One person drives the tool, one guards the red lines below and one keeps the note and speaks for the group. Everyone should question the output.

Before you paste anything

If you would not pin it to a public noticeboard in the building, do not put it into a free AI tool. Free tiers may reuse what you enter for model training, and privacy settings can change without notice. Stop if any of these applies.

- Personal data about an identifiable living person, including student and participant records.
- Special-category data under the UK GDPR, such as health, ethnicity or beliefs, and children's data.
- Anything confidential, unpublished or embargoed, such as a manuscript, a grant draft or a peer review.
- Material whose rights are not yours to give away.

If any applies, choose one of four routes. Anonymise until nobody can be singled out, even by combination. Synthesise a version with the same shape and invented content. Abstract the task and leave out the sensitive specifics. Or abstain, and say why.

Your shared note

Open genai-rt.web.app in any browser. There is no account to create and nothing to install. One device makes the group with a non-identifying name, such as 'Otters', and the passcode the facilitator reads out. It then shares the group name and six-character join code with the others. Choose your track from the dropdown. The note saves as you go. Submit before the lightning round at 13:33 where possible, or during it at the latest. Sharing the note in the public archive is optional.

If the app will not load, use hackmd.io or paper under the headings in your pack, and the facilitator will transcribe the note afterwards.

A personal selection, made in the facilitator's own capacity. It is not the position of the University of Oxford (the facilitator's employer) or of the host, and it is not legal advice.

Run sheet and morning checklist

Copies	1 copy
Sides	Double-sided
Paper	A4 portrait
Finishing	Stapled
Length	5 pages

Kept with the cue cards.

Facilitator run sheet

The quick timeline for the day. It is operational only: the reasoning is in the [facilitator guide](#), and the holdable prompts are on the [cue cards](#). Slide numbers refer to [slides.md](#).

A featured workshop in two 30-minute parts, split by a networking lunch.

Before Part 1 (see the [morning checklist](#))

When	Action
11:45	Projector on, deck open in presenter mode, Wi-Fi and power confirmed.
11:50	Create and pin the fallback-links issue, so that it is there if a group fails over to HackMD later (participants do not post to it).
11:52	Reset the app (genai-rt.web.app): sign in to the private dashboard, clear rehearsal groups, set the session passcode and set the timer's minutes without starting it (see the day-of reset). No pre-named notes are needed.
11:55	Lay two group sheets and one group pack on each table of five, with spare pens and paper. The sheet is the one people read, so put it on top. Post the ten seed ideas around the room, each with its track letter, ready for the 12:26 formation.
11:58	Visible timer ready, title slide up (slide 1). Start at 12:00.

Part 1 · Conceptualising the use of AI in research · 12:00–12:30

Clock	Segment	Slides	Facilitator does
12:00	Welcome and stance	1–5	Start on time. Read the disclaimer. Introduce the critical stance and welcome every level of experience.
12:03	Saved or burned	6	A sentence each with a neighbour, two minutes, then move on.
12:05	A long line of tools	7	AI in the history of tools and the productivity economy: the question is how, with discernment.
12:08	Getting started	8	What these tools are, pick one free tool, how to ask. For all levels.
12:10	Friction is a signal	9	The central idea. Lee et al. and Drosos et al. in a sentence each.
12:14	Spectrum, red lines, ethics	10–12	One breath each. Name the teaching angle (students' data, assessment, disclosure).
12:21	The bigger questions	13	Paradigm shift, fairness, disclosure, posed as live questions for Part 2.
12:23	Human in the loop and the map	14	Accountability stays human. Introduce interwoven and staged oversight.
12:26	What Part 2 asks, then form groups	15–16	Two minutes on the three things Part 2 asks, two on formation. Groups of five form at a seed, with the seed letters posted around the room beforehand so that people walk to a sign. Each group's track follows from its seed, so point it to the ready example for that track in the pack.

Networking lunch · 12:30–13:15

When	Action
12:30	Not your session, but a networking break before Part 2. Leave slide 17 up throughout: it is the only thing telling people to come back to the same table. Re-check the dashboard once (passcode set, no rehearsal groups, timer not running).

Part 2 · Practical AI for research and teaching · 13:15–13:45

Clock	Segment	Slides	Facilitator does
13:15	Settle in and orient	18–24	Re-find groups, agree one problem in 60 seconds and open genai-rt.web.app. Read out the passcode and keep the red lines visible. Point to slides 21–24, but do not present them: the group sheet has the instructions.
13:18	Apply it	25	Press Start countdown as slide 25 goes up. If the settle-in runs late, set Minutes so the countdown still ends at 13:33. Circulate. Stop sensitive data, unstick stalls and offer one provocation to each group. Prioritise a caught error and a clear line for the room. Leave the optional boxes to groups with time left.
about 13:26	Pivot	26	Thirty seconds, keyboards down. Stop building, start interrogating.
13:33	Lightning round	27	Use a visible timer. Give eight groups 45 seconds each. With nine or ten, announce 35 seconds because handovers take time. Jot themes only and leave approvals until after the session. Collect a HackMD link only from a fallback group.
13:40	Close	28–34	Synthesise two or three threads. Take-home actions. Reading. Follow-up. Close at 13:45.

After

When	Action
Same day	Check that every app group is approved on the dashboard, and that any HackMD-fallback group's link is in the issue. Nudge any missing.
+1–2 days	Export the approved, consented work from the dashboard (Export approved (Markdown)), export any HackMD-fallback note too, and archive under <code>submissions/</code> (see <code>submissions/README.md</code>).
+3–5 days	Send the follow-up email with the responsible-use commitment and the archive link.

A personal selection, made in the facilitator's own capacity. It is not the position of the University of Oxford (the facilitator's employer) or of the host, and it is not legal advice.

Morning checklist

Tick these before the room fills. The session is low-tech by design, so most of the risk sits in the first ten minutes.

Technology

- [] Projector or screen working, with the laptop output tested at the room's resolution.
- [] Deck open: either `dist/slides.html` in a browser (presenter view) or `dist/slides.pdf` as a fallback. (Rebuild with `npm run build` if needed.)
- [] Wi-Fi confirmed on a guest device, on the network participants will actually use.
- [] Power: enough sockets and extension leads for tables of five.
- [] `genai-rt.web.app` (the workshop app) opens on a guest device and shows the login form, which is the proof that its code loaded rather than merely the page. The facilitator's Google sign-in to the dashboard works.
- [] A chat tool (ChatGPT, Claude, Gemini or Copilot) reachable on the guest network.
- [] A visible timer (phone, slide timer or kitchen timer) for the segment timings and the 45-second lightning round.

App reset (facilitator)

- [] The workshop app reset for this session: sign in to the private dashboard, clear any rehearsal groups, set the session passcode and set the timer's minutes without starting it. See the [day-of-reset](#) for the full list.
- [] Dashboard reachable and clear of any previous session's submissions (clearing is console-only, so if anything is there, note it and clear it during the lunch reset), ready to approve groups as they submit.

Printed materials (per table of five)

- [] [Group sheet](#), two per table. It holds the whole of Part 2 on one side, so it is the sheet people actually read.
- [] [Group pack](#), one per table: a ready example and first prompts for each track, then the rubric and the fallback note.
- [] [Seed-idea signs](#), ten of them, one per page in A4 landscape, printed and posted around the room before 12:00. Each carries its track letter, and they are generated from the quick group seeds in `project_tracks.md`.
- [] Spare pens and paper (HackMD or paper is the fallback if a group cannot use the app).

Facilitator kit

- [] [Cue cards](#) printed and in order.

- [] Run sheet ready.
- [] Water, a clock you can see and a way to get the room's attention.

Details to confirm

- [] The name, date, venue, times and affiliation in the deck and the emails are yours (they are set to the facilitator's own: Department of Education, University of Oxford, and the University of Westminster on 9 September 2026).
 - [] The placeholders that remain were filled before sending: [VENUE] and [POLL URL] in the pre-workshop email, and [FEEDBACK URL] in the follow-up email.
 - [] The University of Oxford crest is embedded in themes/workshop.css (its source is assets/logo.svg), and the deck was rebuilt after any change.
-

A personal selection, made in the facilitator's own capacity. It is not the position of the University of Oxford (the facilitator's employer) or of the host, and it is not legal advice.

Day-of app reset

Copies	1 copy
Sides	Single-sided
Paper	A4 portrait
Finishing	Loose
Length	1 page

| The passcode, timer and export steps for the lunch break.

App reset checklist for the day

This checklist is operational only. The [facilitator guide](#) holds the rationale, the [run sheet](#) the timeline and the [app README](#) the full deployment steps.

A short reset that gets the live app into a clean, ready state for Westminster on 9 September 2026. The app (`genai-rt.web.app`) is where groups capture their work by default, and HackMD or paper is the fallback.

Before Part 1 (about 11:52, with a re-check during lunch)

- [] Sign in to the facilitator dashboard at `genai-rt.web.app/facilitator.html` with the Google account pbernabeu@gmail.com ('Sign in with Google'). Group devices do not sign in here.
- [] Confirm that sign-in actually completes and that you are on the dashboard, not on the sign-in card. If it bounces, the project's Google sign-in needs checking in the Firebase console, which is a pre-deployment concern to fix before the room fills.
- [] Clear the rehearsal and test groups. Deletion is console-only, with no button on the dashboard. In the Firestore console, delete stray test documents from the `groups` collection and the matching `groupNames` entry, or the name stays reserved.
- [] Confirm that the facilitator group list is clear of rehearsal entries (that is where draft and submitted test groups show) and that the public board at `genai-rt.web.app/dashboard.html` is empty.
- [] Set today's passcode in the Session passcode panel and write it where you can read it out (it overwrites any rehearsal value). The status should then read 'Passcode set. Read it out to the room.' Confirm that the public dashboard opens with this passcode too. Until it is set, no group can start.
- [] In the Session timer panel, set Minutes (the default is 15, the build window) but do not start it yet. Set it to 14 rather than 15, so the countdown ends at 13:32 and leaves a minute to submit inside the build window. Pressing Start again with a new number simply restarts the countdown, so if the settle-in runs late, set Minutes to the time remaining until 13:32.
- [] Have the 'Open the workshop app' slide ready, with the large URL `genai-rt.web.app` and the scannable QR code.
- [] During lunch, glance at the dashboard once more: passcode set, no stray groups, timer not running.

At the start of Part 2

- [] Read out the passcode and the URL (`genai-rt.web.app`), and show the QR slide. 'One device per group creates, then reads out its group name and the six-character code. Everyone else types both and leaves the passcode blank.'
- [] Remind groups to choose a group name that does not identify anyone (for example 'Otters' or 'Team Kelp'), which is the one name they type on the login form.
- [] Press Start countdown at 13:18, when the 'Fifteen minutes, starting now' slide goes up, and leave it unstarted through the 13:15 settle-in. The corner chip is advisory: a calm countdown that never locks their form. Once it expires, it reads 'Time's up' until you reset it.

During

- [] Watch the count line, 'N groups · M awaiting review'.
- [] On each submission: Approve, Reopen... (it prompts for a short note) or Rename... as needed.
- [] Approve only when the group is done. Approval shows the work on the session dashboard and makes the group's form read-only, because owners can edit only while a note is in draft or reopened. The join code is wiped too, so no new device can join, and Reopen is no longer offered.
- [] Rename any identifying name the moment you spot one (Rename..., keeping it short and non-identifying).
- [] If a phase overruns, type a new number of Minutes and press Start again, which restarts the countdown at that length.

After Part 2 (once groups have submitted)

- [] Export approved (Markdown) from the dashboard is the default path. It exports only approved and consented groups and downloads `YYYY-MM-DD_genai-rt-submissions.md`, dated with the day you export.
- [] Check it before committing (names and content), rename it to the workshop date if you exported later, then commit under `submissions/`.
- [] Alternatively, from your own machine if pre-configured, run `npm run archive:pr` (it needs `GENAI_RT_PROJECT` and `GENAI_RT_API_KEY` set and `gh` already logged in). To preview, run `npm run archive:pr -- --dry-run`. If either is not ready, use the Markdown export above.

Troubleshooting

- Sign-in fails: allow pop-ups for the site and try again, since the pop-up path is the reliable one. Then confirm that you are using the pre-registered Google account, and check that Google sign-in is enabled for the project in the Firebase console.
- A group cannot start: the passcode is not set, or they have a typo. Re-read it, since the create path needs the session passcode.
- Reopen... or Rename... does nothing: the browser has suppressed this page's dialogs, which it offers to do after several in a row. Reload the dashboard page.
- Timer not showing: glance at your dashboard's Session timer status line first. It reads 'Running: MM:SS left ...' when live and 'Timer is off.' when stopped. The group chip appears while the countdown is running and reads 'Time's up' after it expires.
- App wobbles: tell the room to switch to HackMD (hackmd.io), with the same headings. In-progress edits are queued on the device and sync on reconnect, so nothing already typed is lost.

A personal selection, made in the facilitator's own capacity. It is not the position of the University of Oxford (the facilitator's employer) or of the host, and it is not legal advice.

Part 4 — put up around the room

Seed-idea signs

Copies	1 copy
Sides	SINGLE-SIDED
Paper	A4 LANDSCAPE
Finishing	Heaviest paper available
Length	10 pages

Posted around the room before 12:00, spread out so ten groups can gather without crowding. Do not let the printer rotate or shrink these to fit portrait.

TRACK

A

Stress-test a study design

TRACK

A

**Critique an assessment or
marking rubric**

TRACK

B

**Turn messy notes into owned
actions**

TRACK

B


Tame an email or admin backlog

TRACK

B


Plan a module, project or paper

TRACK



A one-page explainer of a finding

TRACK



**A quick visualisation or
interactive teaching aid**

TRACK

D

A lay summary or press angle

TRACK

D

**Explain a hard concept for
students**

TRACK

D

A social thread from a paper

Table numbers

Copies	1 copy
Sides	SINGLE-SIDED
Paper	A4 LANDSCAPE
Finishing	Heaviest paper available
Length	10 pages

| One on each table, readable across the room.

1

TABLE

2

TABLE

3

TABLE

4

TABLE

5


TABLE

6

TABLE

7

TABLE



TABLE

9

TABLE

10

TABLE