

Worked examples: one per track

A worked example is a fallback. Start with the seed your group formed around or a member's non-confidential problem. Use the example for your track only when neither is available. It also gives you something to critique on paper if the Wi-Fi fails.

The material below is synthetic and safe to paste. It contains no identifiable people, data or unpublished work, so a group can begin at once. Spend a minute on the data decision aid anyway. Say what you would remove from your own material, then record it under Data Security & Ethics. Each example is small enough to start in the short build window and should lead to a useful caught error.

Track A · Methodological Blind-Spot Detector

Scenario

You are about to submit the study below. Use an AI tool as a critical reviewer, then verify which of its objections are real, which are generic and which are wrong.

Synthetic methods snippet (safe to paste): 'We ran a cross-sectional online survey of 200 marketing managers recruited through a LinkedIn post. Each manager completed a 12-item self-report scale of "authentic leadership" and, in the same questionnaire, rated their own team's "innovativeness" on a 5-item scale. Authentic leadership correlated positively with team innovativeness ($r = .46$, $p < .001$). We conclude that authentic leadership causes higher team innovation and recommend leadership training to raise innovation.'

First move

Ask it openly what is wrong with the design, without naming the lenses, because a prompt that lists the failure modes also hides them. Classify each point it raises as real, generic or wrong. Only then run the second pass across explicit lenses (sampling and recruitment, common-method bias, construct validity, causal inference and generalisability), and note what the guardrails added.

What to watch for

A useful answer should flag the causal claim from cross-sectional data, common-method or single-source bias (the same person rates both variables), self-selection via LinkedIn and self-report of one's own team. Look for vague, always-true objections and invented statistics or citations. Record them as caught errors.

Track B · Accessible Executive-Function Layer

Scenario

Turn a messy set of meeting notes into a clear action list, then check that the tool did not quietly invent or misattribute anything.

Synthetic meeting notes (safe to paste; roles only, no real people): 'Centre meeting. Talked about the seminar series. Someone should sort speakers for the autumn, maybe the deputy director? Budget underspend, need to use it before July or we lose it. The new PhD reps raised workload worries. Website is out of date (old staff list). Agreed to revisit the mentoring scheme but no one said who. Ethics turnaround is too slow. Chase the committee. Next meeting in three weeks, same room hopefully.'

First move

Ask for a structured action list with owner, deadline and priority. Ask it to place unclear items, or items without an owner, in a separate list. Turn the resulting prompt into a template your group could use again.

What to watch for

Does the tool invent owners or dates that the notes never stated ('Deputy Director, by 30 June')? Does it silently drop items with no owner? Does it over-formalise a vague discussion? Use a safeguard such as 'every action must trace to a line in the notes; no invented owners'.

Track C · Rapid Prototyping for Knowledge Translation (*technical stretch*)

Scenario

Build a tiny interactive that communicates a finding, then verify it before sharing it.

Synthetic finding and data (safe to paste): 'In a (made-up) study, average reading time for a one-page brief fell as font size rose, then rose again when the font got too large. Synthetic data, mean seconds: 10 pt → 95 s; 12 pt → 78 s; 14 pt → 70 s; 16 pt → 72 s; 18 pt → 81 s.'

First move

Ask a tool with a preview (for example, Claude artifacts or Gemini's Canvas) to build a small, self-contained HTML page with a labelled bar or line chart of these values, plus a one-sentence plain-language takeaway. Keep a verification log as you go.

What to watch for

Do the plotted numbers match the data exactly? Is the takeaway honest about the U-shape (fastest reading around 14 pt), or does it over-simplify to 'bigger is better'? Check colour contrast and that the axes are labelled. Then add a sixth point of your own, 20 pt, and see whether the axis, the scale and the takeaway update correctly or whether the tool hard-coded them. The fact that it rendered does not mean that it is correct, so log every gap.

Track D · Public Engagement Translator

Scenario

Translate a hedged, cautious finding for three audiences, then audit each version for accuracy drift.

Synthetic abstract (safe to paste): 'In a small, preliminary, correlational study of 60 undergraduates, students who reported using a structured note-taking app tended to report slightly higher exam confidence (not exam marks). The effect was modest, the sample was not representative, and no causal claim can be made. Replication is needed.'

First move

Produce three versions (an interested public, a policy reader and secondary-school pupils), then line each up against the source and mark every over-claim, dropped caveat or false certainty.

What to watch for

Watch the tool quietly upgrade 'tended to report slightly higher confidence' into 'boosts exam results', drop 'small, preliminary, correlational' or add a headline the evidence cannot bear. Create a fidelity checklist: every claim is traceable to the source, all caveats survive and no causal language appears.

Disclaimer: these materials and the tools they reference are the facilitator's own choices, made in a personal capacity, and do not represent the views or official position of the University of Oxford, the facilitator's employer, or of the host, the University of Westminster. The University of Oxford accepts no liability for the selection, use or outcomes of any third-party tool. Participants remain solely responsible for compliance with their own institutional policy, the UK GDPR and research ethics. This wording is a template and is not legal advice.

Fallback note and rubric

Use this page if the app fails or when your group wants to record a fuller reflection after the session. It does not need attention during the fifteen-minute build.

During the build, use the app to record three things: the result of using a tool on your problem, one or two caught errors and the line you will give the room at 13:33. That line may name a limitation, a checkpoint you would require before use or how your field over- or under-uses AI for this task. The material below supports the fallback and the debrief.

If the app will not work

Open a note at hackmd.io, or take the paper sheet from the facilitator, and copy the headings below into it. Make rough notes as you work. They preserve what happened better than a tidy account written at the end. Give the facilitator the link or sheet before the lightning round at 13:33 where possible, or during it at the latest.

If HackMD will not let you sign up, switch to paper at once. The headings work just as well there, and a sign-up is never worth the minutes that troubleshooting it costs.

Group ____ · Track A / B / C / D / Own problem · 9 September 2026

The problem:

The artefact, and the tool that made it:

Errors caught (what happened, how we caught it, where a human must stay):

Our line for the lightning round:

Optional, if minutes remain

Automation-steering map (the phases, what the tool ran, where we steered):

Oversight interwoven or staged, and why:

Field reflection (over- or under-used here, and what we would do about it):

Archive this note publicly? yes / no

The rubric after the session

For a short applied session, 3 is a reasonable result. A sharp caught error tells you more than an unexamined 5. Use 2 and 4 for work that falls between two anchors.

Dimension	1 Nascent	3 Developing	5 Robust
Project definition. Is the problem, and how you would know it helped, clear?	a vague or toy task, with no success criterion	a problem named, with a rough sense of better	bounded and owned, with a criterion you could test
Technology stack. Is this the right tool on the spectrum, and why?	one tool by default, with alternatives unconsidered	the choice explained, with some sense of the spectrum	placed deliberately, weighing control, data exposure and effort
Data security and ethics. Were the red lines held?	sensitive material pasted without thought	red lines mostly held and the data de-identified	red lines held, with disclosure, IP, consent and fairness reasoned
Financial and scalability. Does the free tier hold?	no thought given to cost or scale	limits noted, with a rough sense of what scaling needs	clear on lock-in and on what breaks first, cost, limits or trust
Human in the loop. Are there checkpoints, verification and accountability?	the output trusted as it stands	some verification, and one or two checkpoints named	oversight chosen and its weakness offset, with one named owner

To record the scores, add a line to your note for each dimension with its number and a sentence saying why.

Worth a few minutes afterwards

Say what thinking you still had to do once the tool had done the writing, and what would go if you handed that over as well. Note who gains from working this way and who is left behind, whether by first or second language, support, career stage or background. Set down whether you would disclose the use of AI here, and whether you would also disclose the conditions it offset. Then judge as honestly as you can whether your field is over- or under-using AI for this kind of task, and what you would want done about it over the next two years.

Disclaimer

These materials and the tools they reference are the facilitator's own choices, made in a personal capacity. They do not represent the views or official position of the University of Oxford, the facilitator's employer, or of the host, the University of Westminster. The University of Oxford accepts no liability for the selection, use or outcomes of any third-party tool. Participants remain solely responsible for compliance with their own institutional policy, the UK GDPR and research ethics.

This wording is a template and is not legal advice.